

# Performance

- What does it mean when you say “system A is of a higher performance as compared to system B”?

# Performance

- What does it mean when you say “system A is of a higher performance as compared to system B”?
- “Time taken to perform a certain task T on system A is less than that taken to perform the same task T on system B.”
  - Note that this statement is specific to the task T. For some other task T', it may be possible that system B is the better performer.

# Performance

|  |   |
|--|---|
| <pre>int f = 1; for(int i=2; i&lt;=n; i++) {     f = f * i; } cout &lt;&lt; f;</pre> | <pre>int f = 1; for(int i=2; i&lt;=2*n; i++) {     if (i &lt;= n)         f = f * i; } cout &lt;&lt; f;</pre> |
| Program A  | Program B   |

- System A: Program A is run on processor type P
- System B: Program B is run on processor type P
- Which system displays higher performance?

# Performance

```
.main:
[0]      mov r2, 3
[4]      mov r1, 1
.loop:
[8]      mul r1, r1, r2
[12]     sub r2, r2, 1
[16]     cmp r2, 1
[20]     bgt .loop
[24]     st r1, 4[r0]
[28]     end
```

How many instructions are executed?

# Performance

```
.main:
[0]      mov r2, 3
[4]      mov r1, 1
.loop:
[8]      mul r1, r1, r2
[12]     sub r2, r2, 1
[16]     cmp r2, 1
[20]     bgt .loop
[24]     st r1, 4[r0]
[28]     end
```

- Number of static instructions = 8
- Number of dynamic instructions = 12

# Performance

|  |   |
|--|---|
| <pre>int f = 1; for(int i=2; i&lt;=n; i++) {     f = f * i; } cout &lt;&lt; f;</pre> | <pre>int f = 1; for(int i=2; i&lt;=2*n; i++) {     if (i &lt;= n)         f = f * i; } cout &lt;&lt; f;</pre> |
| Program A  | Program B   |

- System A: Program A is run on processor type P
- System B: Program B is run on processor type P
- System A displays higher performance because program A has fewer dynamic instructions (assuming every instruction takes equal amount of time)

# SimpleRISC Processor Frequency

- How long does it take to execute a single instruction on our SimpleRISC processor?

# SimpleRISC Processor Frequency

- How long does it take to execute a single instruction on our SimpleRISC processor?
- Work out the longest path in the combinational circuitry. Processor clock period should be greater than the time taken to execute this longest path. Let this clock period be  $t$  seconds. Processor frequency  $f = 1/t$ .
- Every instruction is executed in  $t$  s.
- Every instruction is executed in 1 cycle in our SimpleRISC processor.
  - Cycles per instruction,  $CPI = 1$
  - Instructions per cycle,  $IPC = 1$



# Computing the Time a Program Takes

$$\begin{aligned}\tau &= \#seconds \\ &= \frac{\#seconds}{\#cycles} * \frac{\#cycles}{\#instructions} * (\#instructions) \\ &= \underbrace{\frac{\#seconds}{\#cycles}}_{1/f} + \underbrace{\frac{\#cycles}{\#instructions}}_{CPI} * (\#instructions) \\ &= \frac{CPI * \#insts}{f}\end{aligned}$$

- \* **CPI** → **Cycles** per instruction
- \* **f** → **frequency** (cycles per second)

# The Performance Equation

$$P \propto \frac{IPC * f}{\#insts}$$

- \* **IPC** → 1/CPI (Instructions per Cycle)
- \* What are the **units** of performance ?
  - \* **ANSWER** : arbitrary

# Number of Instructions (#insts)

**Static Instruction:** The **binary** or executable of a **program**, contains a list of *static instructions*.

**Dynamic Instruction:** A **dynamic instruction** is a running instance of a static instruction, which is created by the **processor** when an instruction enters the pipeline.

- \* Note that these are **dynamic** instructions
  - \* **NOT** static instructions
- \* A smart compiler can reduce the number of executed instructions

# Number of Instructions(#insts) – 2

- \* Dead code removal

- \* Often **programmers** write **code** that does not determine the final output
- \* This code is **redundant**
- \* It can be identified and removed by the **compiler**

- \* Function inlining


- \* Very small **functions** have a lot of **overhead** → call, ret instructions, register spilling, and restoring
- \* Paste the code of the **callee** in the code of the **caller** (known as **inlining**)



# Computer Organisation and Architecture

Smruti Ranjan Sarangi,  
IIT Delhi

## Chapter 9 Principles of Pipelining

Also available as  Book

# Computer Organisation and Architecture


Smruti Ranjan Sarangi



These slides are meant to be used along with the book: Computer Organisation and Architecture, Smruti Ranjan Sarangi, McGrawHill 2015  
Visit: <http://www.cse.iitd.ernet.in/~srsarangi/archbooksoft.html>

# Outline

- \* Overview of Pipelining
- \* A Pipelined Data Path
- \* Pipeline Hazards
- \* Pipeline with Interlocks
- \* Forwarding
- \* Performance Metrics
- \* Interrupts/ Exceptions



# Up till now ....

- \* We have designed a **processor** that can execute all the **SimpleRisc** Instructions
- \* We have look at two **styles** :
  - \* With a **hardwired** control unit
  - \* **Microprogrammed** control unit
    - \* Microprogrammed data path
    - \* Microassembly Language
    - \* Microinstructions




# Designing Efficient Processors

- \* **Microprogrammed** processors are much slower than **hardwired** processors
- \* Even **hardwired** processors
  - \* Have a lot of **waste** !!!
  - \* We have 5 stages.
  - \* What is the IF stage doing, when the MA stage is active ?
    - \* **ANSWER** : It is **idling**

# The Notion of Pipelining

- \* Let us go back to the **car assembly line**
  - \* Is the **engine shop** idle, when the **paint shop** is painting a car ?
  - \* **NO** : It is building the engine of another car
  - \* When this engine goes to the body shop, it builds the engine of another car, and **so on ....**
- \* **Insight** :
  - \* **Multiple cars** are built at the same time.
  - \* A car **proceeds** from one stage to the next

# Pipelined Processors



- \* The IF, ID, EX, MA, and RW stages process 5 instructions simultaneously
- \* Each instruction proceeds from one stage to the next
- \* This is known as **pipelining**

# Advantages of Pipelining

- \* We keep all parts of the data path, busy all the time
- \* Let us assume that all the 5 stages do the same amount of **work**
  - \* **Without** pipelining, every **T** seconds, an instruction completes its **execution**
  - \* **With** pipelining, every **T/5** seconds, a new instruction completes its **execution**

# Design of a Pipeline


## \* Splitting the Data Path

- \* We divide the data path into 5 parts : IF, OF, EX, MA, and RW

## \* Timing

- \* We insert latches (registers) between consecutive stages
- \* 4 Latches → IF-OF, OF-EX, EX-MA, and MA-RW
- \* At the negative edge of a clock, an instruction moves from one stage to the next

# Pipelined Data Path with Latches




- \* Add a latch between subsequent stages.
  - \* Triggered by a negative clock edge

# The Instruction Packet

- \* What travels **between stages** ?
  - \* **ANSWER** : the instruction packet
- \* Instruction Packet
  - \* **Instruction contents**
  - \* Program counter
  - \* **All intermediate results**
  - \* Control signals
- \* Every instruction moves with its entire state, no interference between instructions

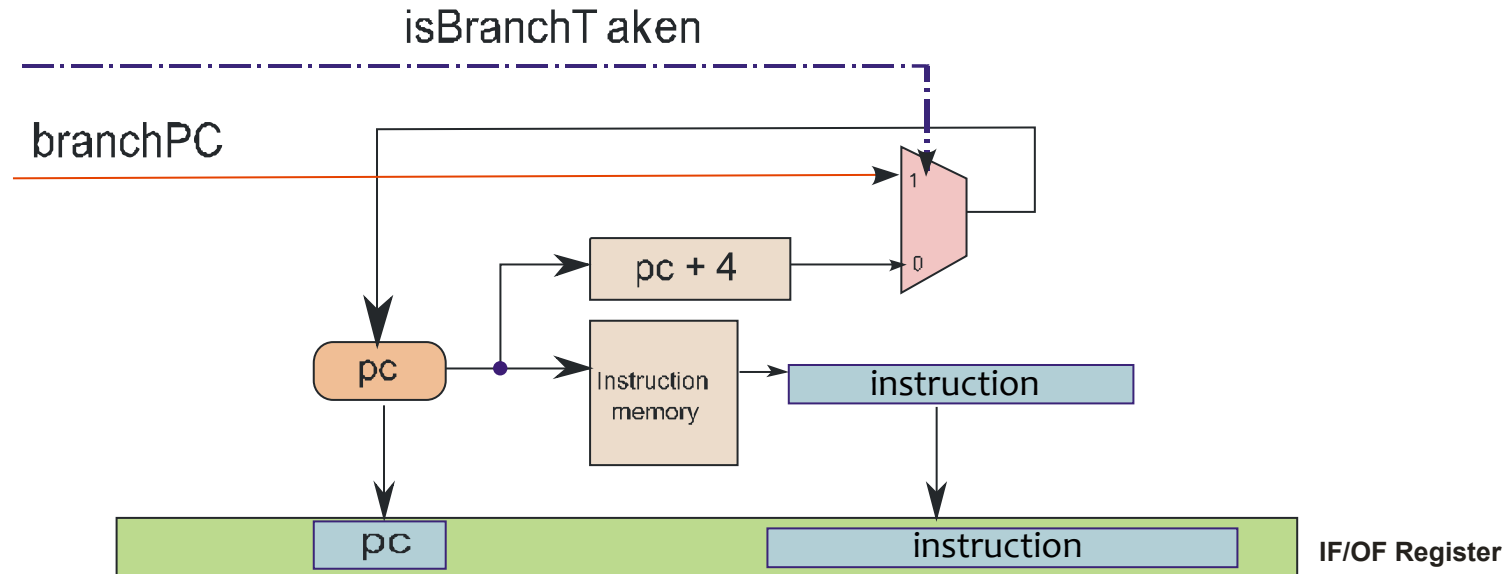
# Outline

- \* Overview of Pipelining
- \* A Pipelined Data Path
- \* Pipeline Hazards
- \* Pipeline with Interlocks
- \* Forwarding
- \* Performance Metrics
- \* Interrupts/ Exceptions



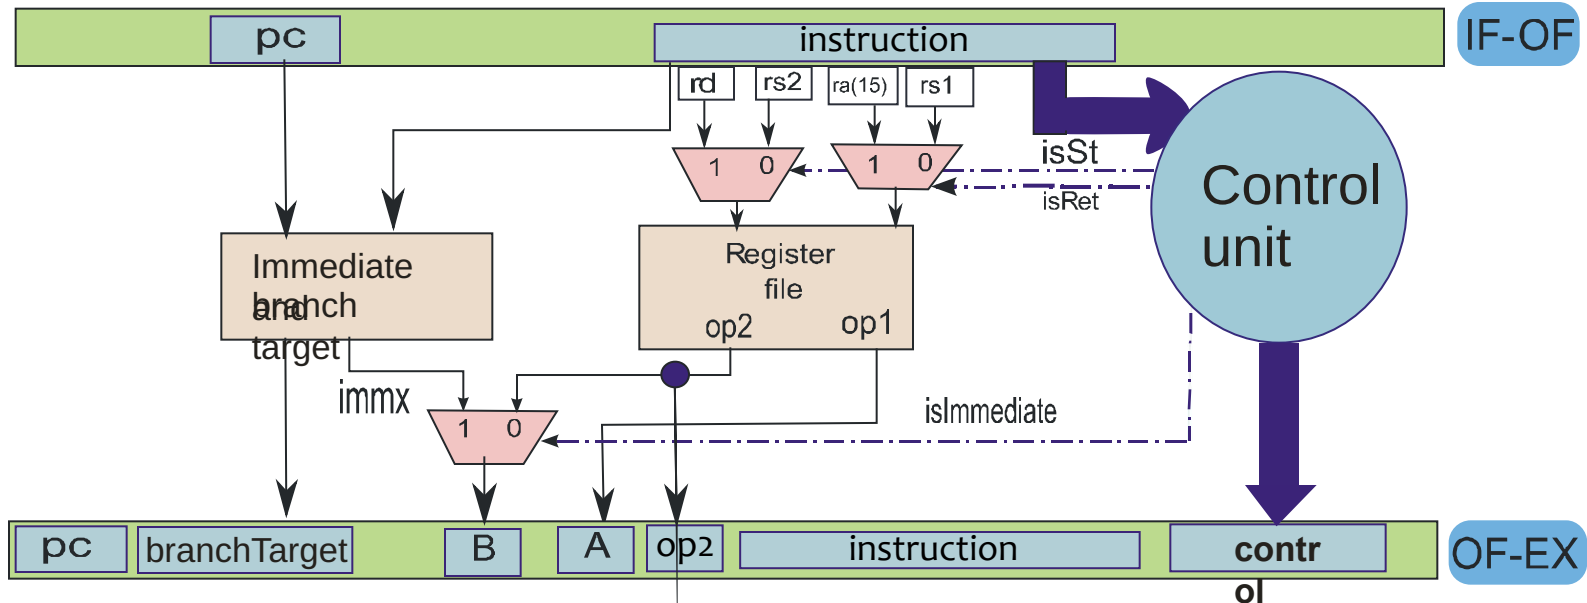


# IF Stage



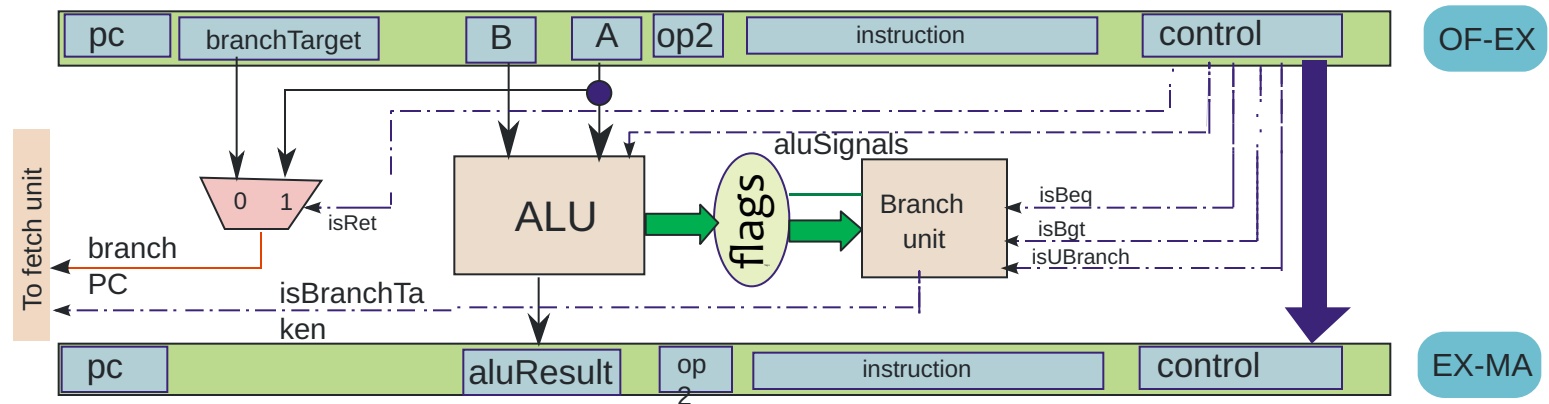
- \* Instruction contents saved in the **instruction** field

# OF Stage



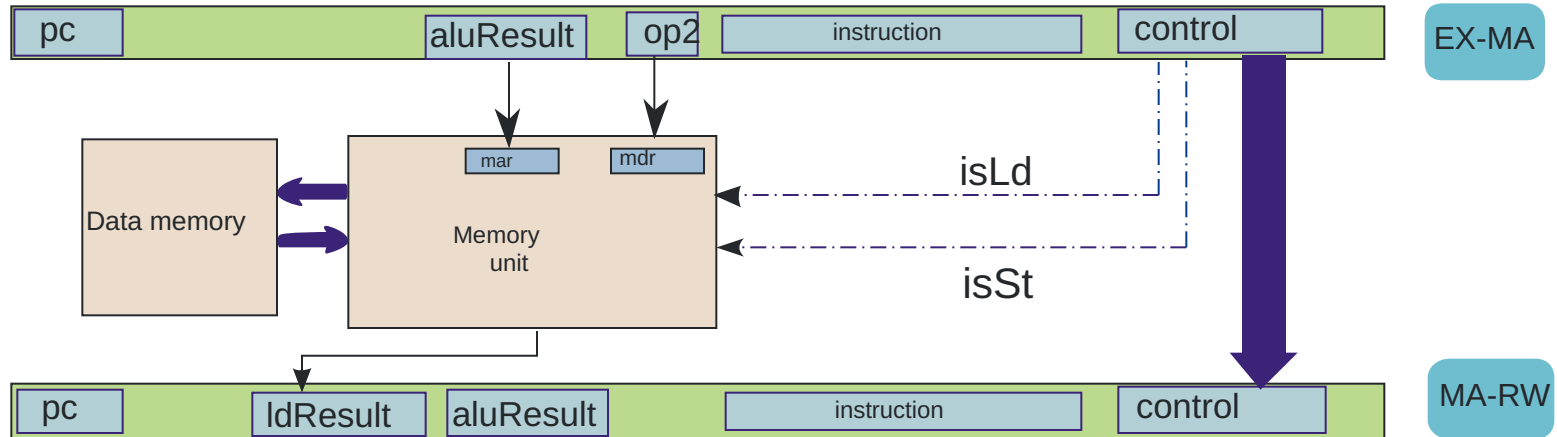
\* **A, B** → ALU Operands, **op2** (store operand),  
control (set of all control signals)

# EX Stage



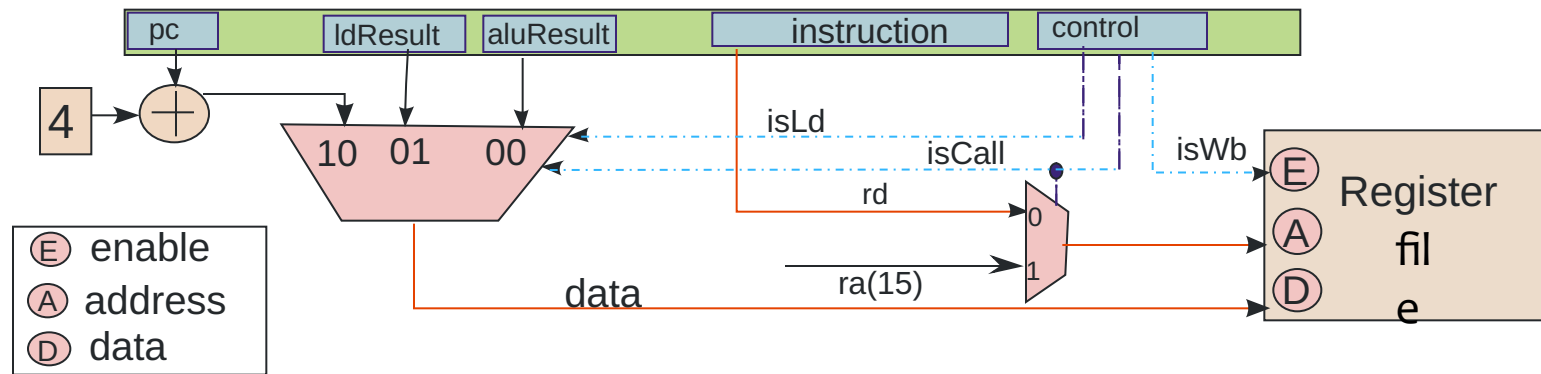
- \* **aluResult** → result of the ALU Operation
- \* **op<sub>2</sub>, control, pc, instruction** (passed from OF-EX)

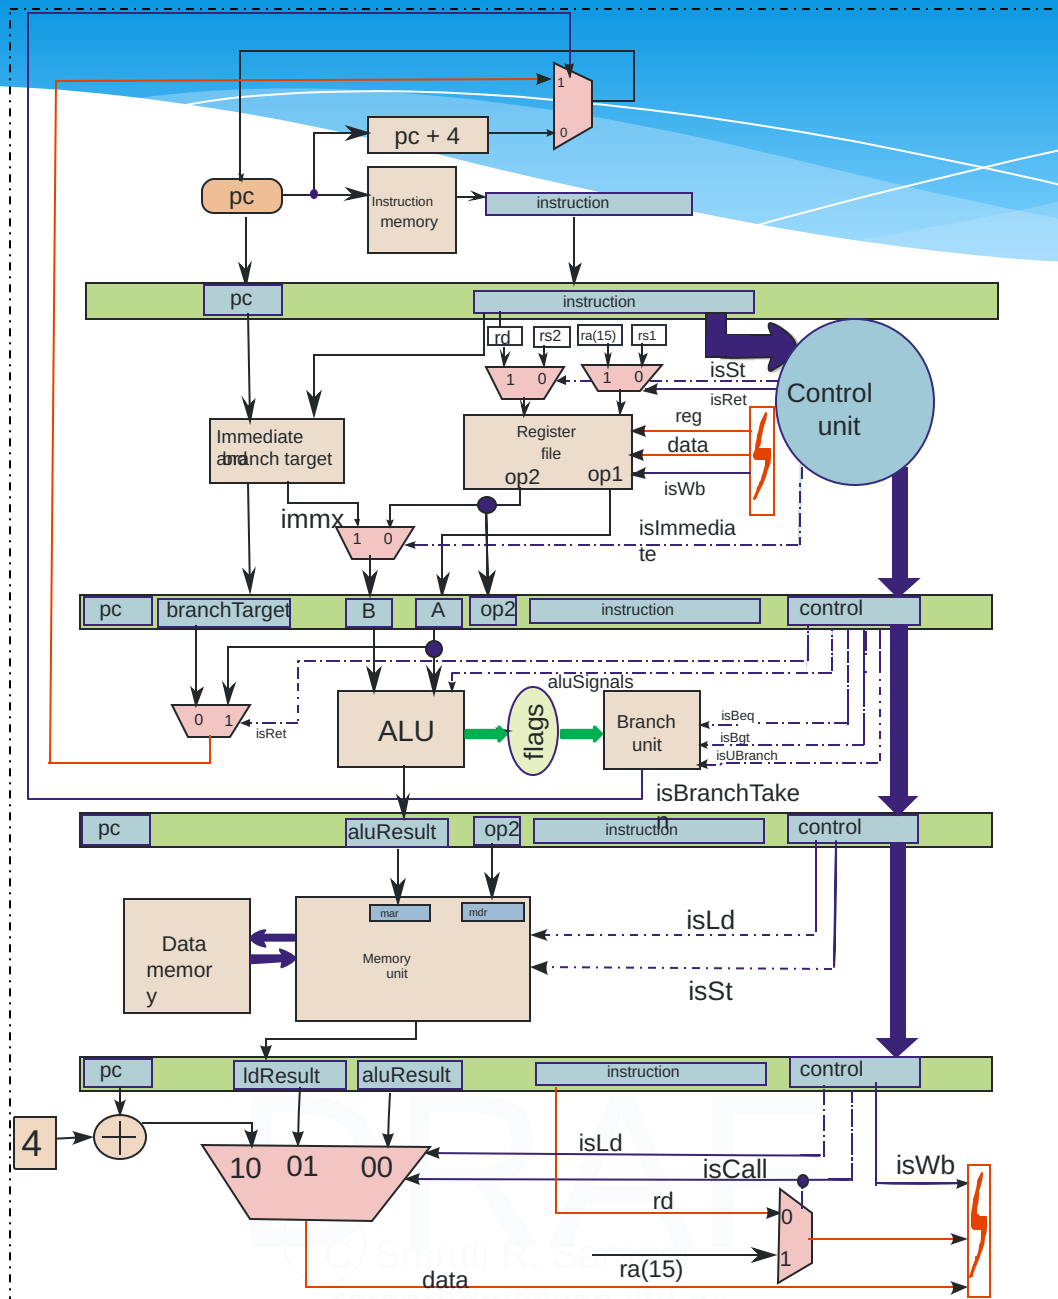
# MA Stage



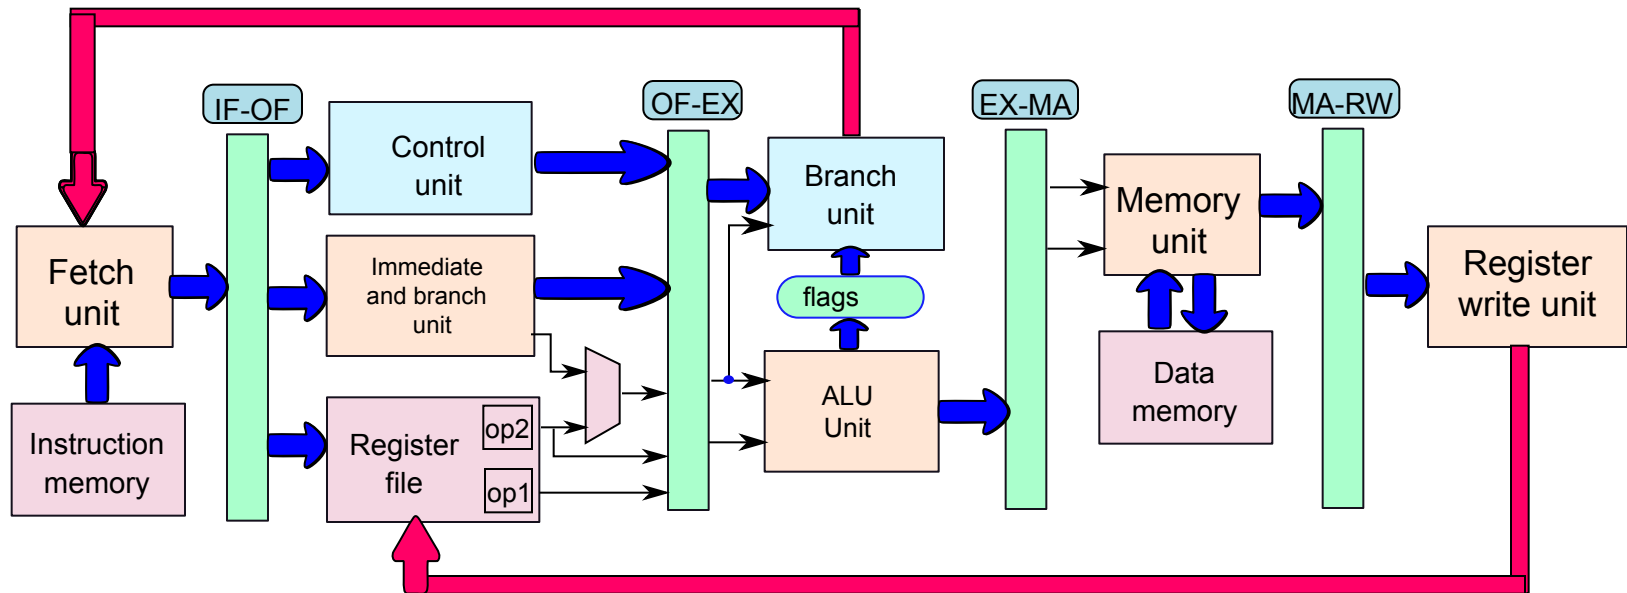
- \* **IdResult** → result of the load operation
- \* **aluResult, control, pc, instruction** (passed from EX-MA)

# RW Stage






# Abridged Diagram



# Outline

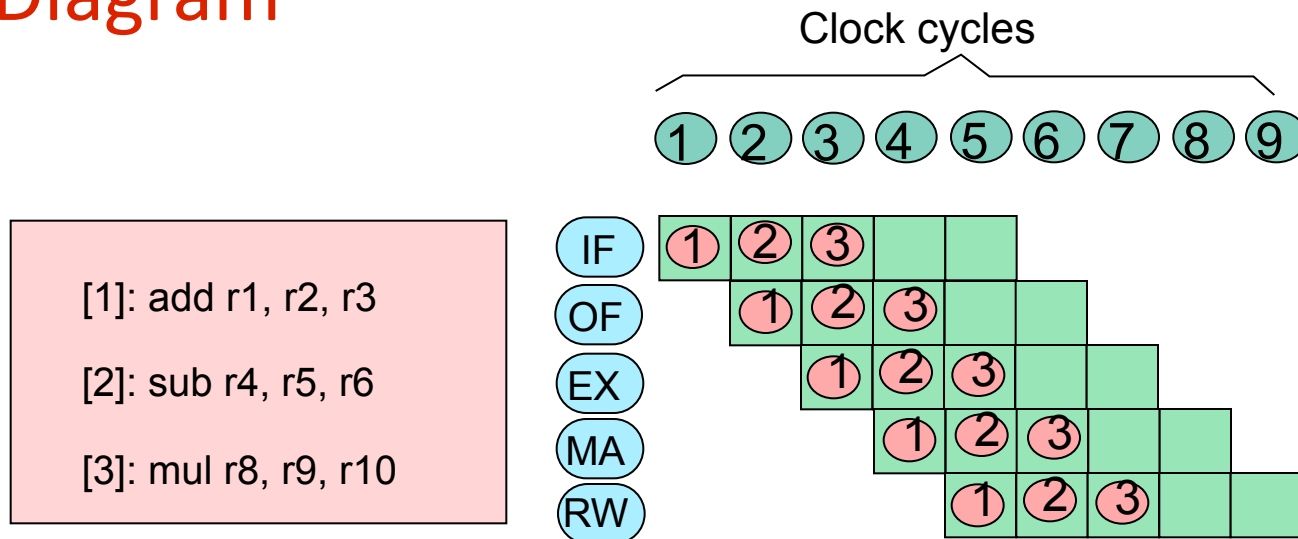
- \* Overview of Pipelining
- \* A Pipelined Data Path
- \* Pipeline Hazards
- \* Pipeline with Interlocks
- \* Forwarding
- \* Performance Metrics
- \* Interrupts/ Exceptions





# Pipeline Hazards

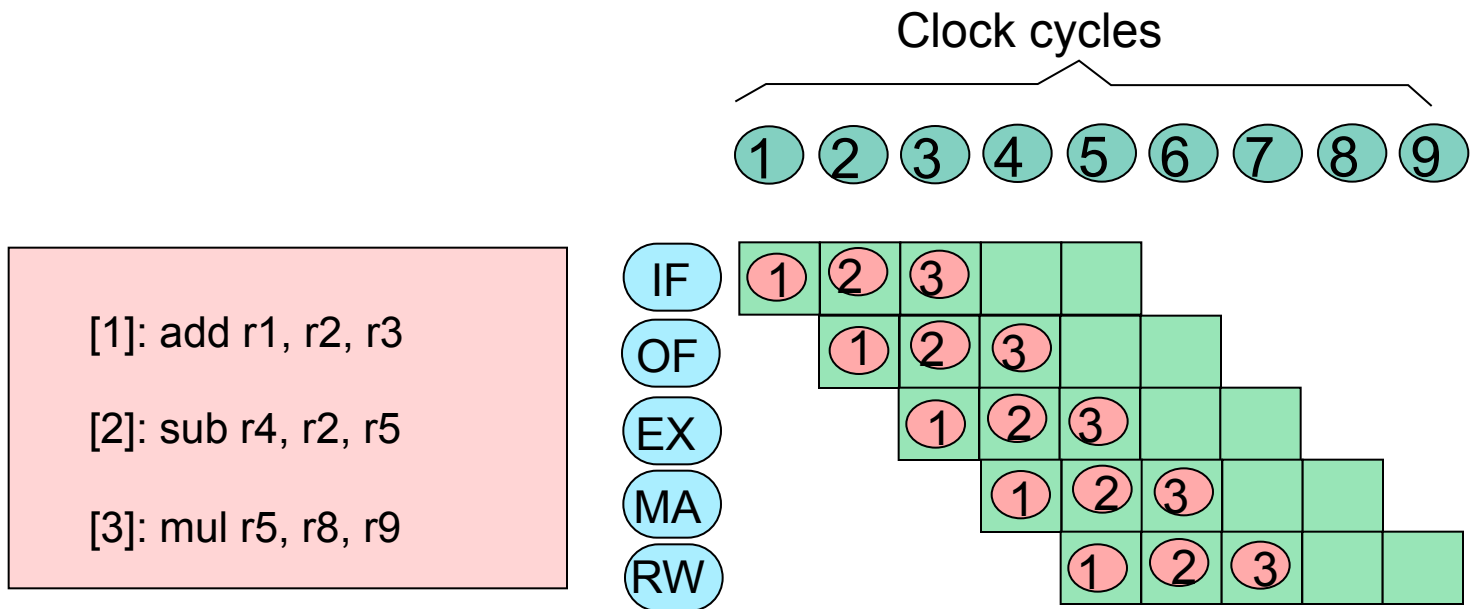
- \* Now, let us consider **correctness**
- \* Let us introduce a new tool → **Pipeline Diagram**



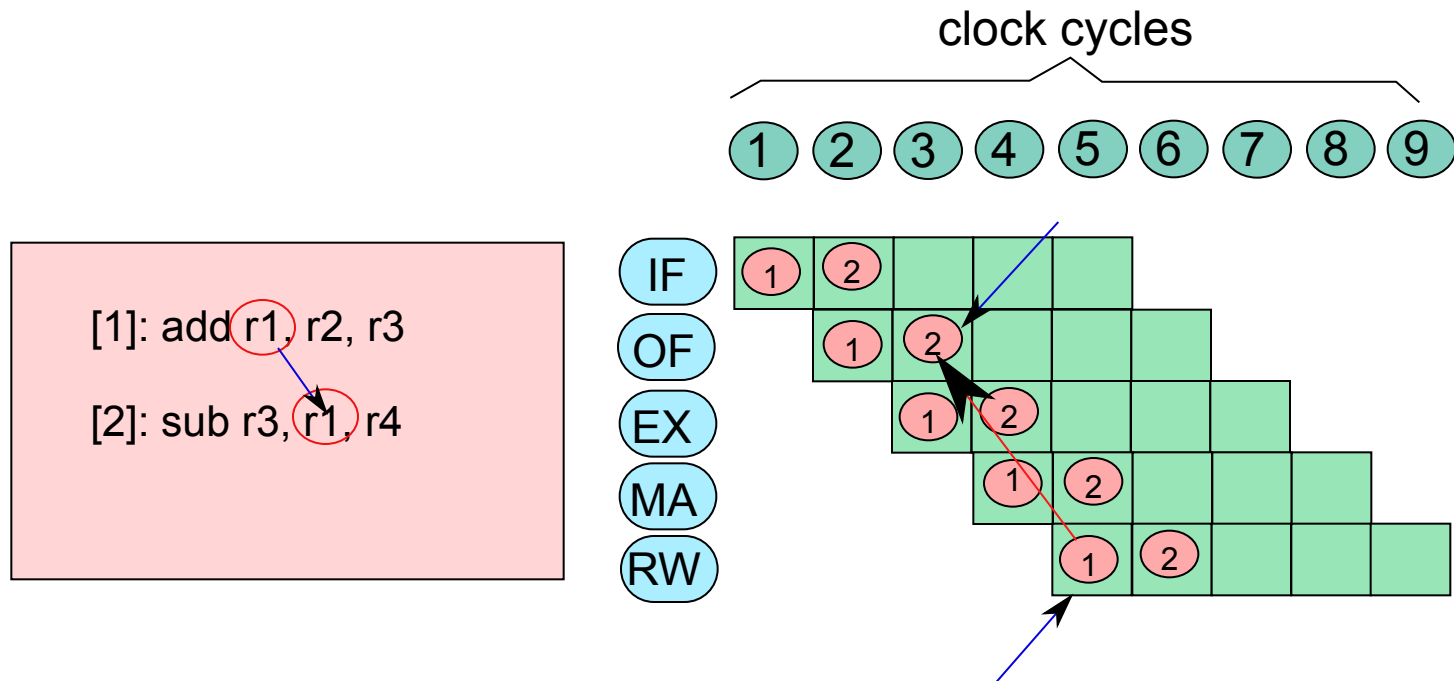
# Rules for Constructing a Pipeline Diagram

- \* It has 5 **rows**
  - \* One per each stage
  - \* The rows are **named** : IF, OF, EX, MA, and RW
- \* Each **column** represents a clock cycle
- \* Each **cell** represents the execution of an instruction in a stage
  - \* It is **annotated** with the name(label) of the instruction
- \* Instructions proceed from one stage to the next across clock cycles

# Example



# Data Hazards



\* Instruction 2 will read incorrect values !!!

# Data Hazard

**Definition:** A **hazard** is defined as the possibility of erroneous execution of an instruction in a pipeline. A data hazard represents the possibility of erroneous execution because of the unavailability of data, or the availability of **incorrect** data.

- \* This situation represents a **data hazard**
- \* In specific,
  - \* it is a **RAW** (read after write) hazard
- \* The earliest we can dispatch instruction 2, is cycle 5

# Other Types of Data Hazards

- \* Our pipeline is in-order

**Definition:** In an in-order pipeline (such as ours), a preceding instruction is always ahead of a succeeding instruction in the pipeline. Modern processors however use out-of-order pipelines that break this rule. It is possible for later instructions to execute before earlier instructions.

- \* We will only have RAW hazards in our pipeline.
- \* Out-of-order pipelines can have WAR and WAW hazards

# WAW Hazards

```
[1]: add r1, r2, r3  
[2]: sub r1, r4, r3
```

- \* Instruction [2] cannot write the value of r1, before instruction [1] writes to it, will lead to a **WAW** hazard

# WAR Hazards

```
[1]: add r1, r2, r3  
[2]: add r2, r5, r6
```

- \* **Instruction** [2] cannot **write** the value of r2, before **instruction** [1] reads it → will lead to a **WAR** hazard



[1] add r1, r2, r3  
[2] ld r4, 8[r1]  
[3] sub r5, r5, r2  
[4] st r4, 12[r2]  
[5] sub r6, r4, r3  
[6] add r7, r4, r3

Q. How many dependencies?

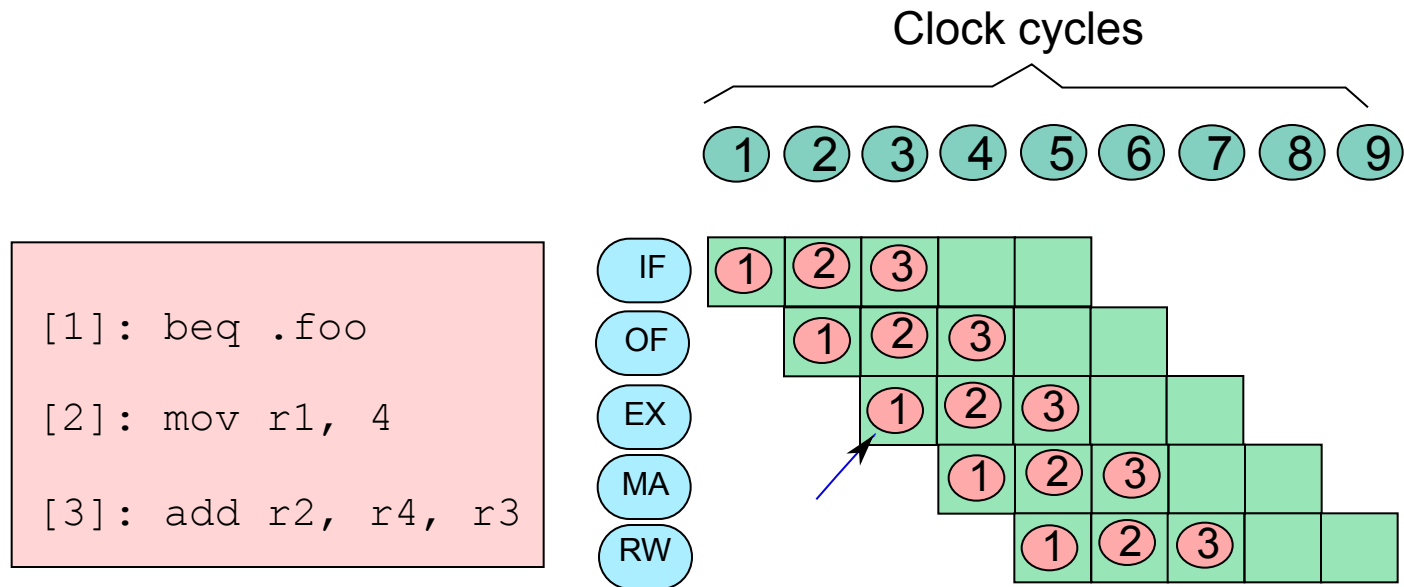
Q. How many dependencies require conflict resolution?

# Control Hazards

```
[1]: beq .foo
[2]: mov r1, 4
[3]: add r2, r4, r3
...
...
.foo:
[100]: add r4, r1, r2
```

- \* If the branch is **taken**, instructions [2] and [3], might get fetched, incorrectly

# Control Hazard – Pipeline Diagram



- \* The two **instructions** fetched immediately after a branch instruction might have been fetched **incorrectly**.

```
[1] mov r1, 0
[2] mov r2, 0
[3] cmp r1, r2
[4] beq .foo
[5] sub r6, r4, r3
[6] add r7, r4, r3
[7] mul r8, r4, r3
...
...
.foo
[100] add r5, r4, r1
```

# Control Hazards

- \* The two **instructions** fetched immediately after a branch instruction might have been fetched **incorrectly**.
- \* These instructions are said to be on the **wrong path**
- \* A **control hazard** represents the **possibility** of **erroneous** execution in a **pipeline** because instructions in the **wrong path** of a **branch** can possibly get **executed** and save their results in memory, or in the register file

# Structural Hazards

- \* A **structural hazard** may occur when two instructions have a conflict on the same set of resources in a cycle
- \* Example :
  - \* Assume that we have an add instruction that can read one operand from memory
  - \* `add r1, r2, 10[r3]`

# Structural Hazards - II

```
[1]: st r4, 20[r5]  
[2]: sub r8, r9, r10  
[3]: add r1, r2, 10[r3]
```


- \* This code will have a structural hazard
  - \* [3] tries to read 10[r3] (MA unit) in cycle 4
  - \* [1] tries to write to 20[r5] (MA unit) in cycle 4
- \* Does not happen in our pipeline

# Solutions in Software

- \* Data hazards

- \* Insert **nop** instructions, reorder **code**

```
[1]: add r1, r2, r3  
[2]: sub r3, r1, r4
```



```
[1]: add r1, r2, r3  
[2]: nop  
[3]: nop  
[4]: nop  
[5]: sub r3, r1, r4
```




# Code Reordering

```
add r1, r2, r3
add r4, r1, 3
add r8, r5, r6
add r9, r8, r5
add r10, r11, r12
add r13, r10, 2
```

# Code Reordering

```
add r1, r2, r3
add r4, r1, 3
add r8, r5, r6
add r9, r8, r5
add r10, r11, r12
add r13, r10, 2
```




```
add r1, r2, r3
add r8, r5, r6
add r10, r11, r12
nop
add r4, r1, 3
add r9, r8, r5
add r13, r10, 2
```

# Control Hazards

- \* **Trivial Solution** : Add two nop instructions after every branch
- \* **Better solution** :
  - \* Assume that the two instructions fetched after a branch are **valid** instructions
  - \* These instructions are said to be in the **delay slots**
  - \* Such a branch is known as a **delayed branch**

# Example with 2 Delay Slots

```
add r1, r2, r3
add r4, r5, r6
b .foo
add r8, r9, r10
```




```
b .foo
add r1, r2, r3
add r4, r5, r6
add r8, r9, r10
```

- \* The **compiler** transfers **instructions** before the branch to the **delay slots**.
- \* If it cannot find 2 valid **instructions**, it inserts **nops**.

# Outline

- \* Overview of Pipelining
- \* A Pipelined Data Path
- \* Pipeline Hazards
- \* Pipeline with Interlocks
- \* Forwarding
- \* Performance Metrics
- \* Interrupts/ Exceptions



# Why interlocks ?

- \* We cannot always **trust** the **compiler** to do a good job, or even introduce **nop** instructions correctly.
- \* **Compilers** now need to be tailored to specific hardware.
- \* We should ideally not expose the details of the **pipeline** to the **compiler** (might be confidential also)
- \* Hardware mechanism to enforce correctness → **interlock**

# Two kinds of Interlocks

- \* Data-Lock

- \* Do not allow a consumer instruction to move beyond the OF stage till it has read the correct values. Implication : Stall the IF and OF stages.

- \* Branch-Lock

- \* We never execute instructions in the wrong path.
- \* The hardware needs to ensure both these conditions.

# Comparison between Software and Hardware

| Attribute   | Software  | Hardware(withinterlocks)  |
|-------------|---|---|
| Portability | Limited to a specific processor                               | Programs can be run on any processor irrespective of the nature of the pipeline                                   |
| Branches    | Possible to have no performance penalty, by using delay slots | Need to stall the pipeline for 2 cycles in our design   |
| RAW hazards | Possible to eliminate them through code scheduling            | Need to stall the pipeline  |
| Performance | Highly dependent on the nature of the program                 | The basic version of a pipeline with interlocks is expected to be slower than the version that relies on software |



# Conceptual Look at Pipeline with Interlocks

```
[1]: add r1, r2, r3  
[2]: sub r4, r1, r2
```

- \* We have a **RAW hazard**
- \* We need to **stall**, instruction [2] at the OF stage for 3 cycles.
- \* We need to keep sending **nop** instructions to the **EX stage** during these 3 cycles

# Example




bubble




[1]: add r1, r2, r3

[2]: sub r4, r1, r2

Clock cycles



# Example



# A Pipeline Bubble


- \* A **pipeline bubble** is inserted into a stage, when the **previous stage** needs to be stalled
- \* It is a **nop** instruction
- \* To insert a **bubble**
  - \* Create a **nop instruction packet**
  - \* OR, Mark a designated **bubble bit** to 1

# Bubbles in the Case of a Branch Instruction



bubble

```
[1]: beq. foo
[2]: add r1, r2, r3
[3]: sub r4, r5, r6
....
....
.foo:
[4]: add r8, r9, r10
```



# Control Hazards and Bubbles

- \* We know that an instruction is a **branch** in the OF stage
- \* When it reaches the EX stage and the branch is taken, let us **convert** the instructions in the IF, and OF stages to **bubbles**
- \* Ensures the **branch-lock** condition

IF

IF-OF Latch

Inst\_OF

OF

OF-EX Latch

Inst\_EX

EX

EX-MA Latch

Inst\_MA

MA

MA-RW Latch

Inst\_RW

RW

# Ensuring the Data-Lock Condition

- \* When an **instruction** reaches the **OF** stage, check if it has a **conflict** with any of the instructions in the **EX, MA, and RW** stages
- \* If there is **no conflict**, **nothing** needs to be done
- \* Otherwise, **stall** the **pipeline** (IF and OF stages only)



### Algorithm 5: Algorithm to detect conflicts between instructions

**Data:** instructions, [A], and [B]

**Result:** conflict exists (**true**), no conflict (**false**)

```
if [A].opcode  $\in$  (nop,b,beq,bgt,call) then
    /* Does not read from any register */
    return false
end
if [B].opcode  $\in$  (nop, cmp, st, b, beq, bgt, ret) then
    /* Does not write to any register */
    return false
end
/* Set the sources */
src1  $\leftarrow$  [A].rs1
src2  $\leftarrow$  [A].rs2
if [A].opcode = st then
    src2  $\leftarrow$  [A].rd
end
if [A].opcode = ret then
    src1  $\leftarrow$  ra
end
hasSrc1  $\leftarrow$  true
if ([A]  $\in$  (not, mov)) hasSrc1  $\leftarrow$  false
```

```

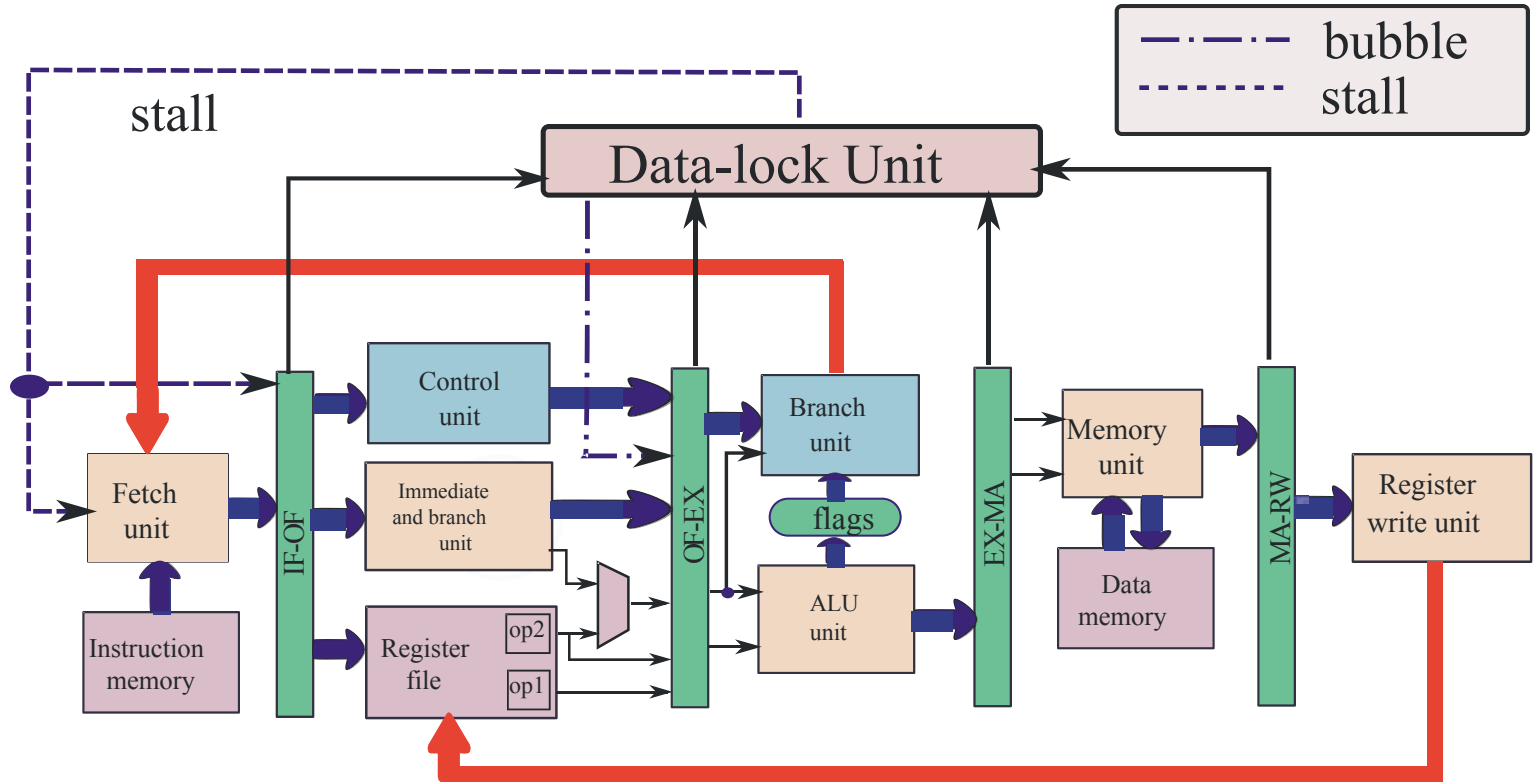
dest ← [B].rd
if [B].opcode = call then
    dest ← ra
end
/* Check the second operand to see if it is a register */
hasSrc2 ← true
if [A].opcode ≠ (st) then
    if [A].I = 1 then
        hasSrc2 ← false
    end
end
/* Detect conflicts */
if (hasSrc1 = true) and (src1 = dest) then
    return true
end
else if (hasSrc2 = true) and (src2 = dest) then
    return true
end
return false

```

# How to Stall a Pipeline ?

- \* Disable the **write functionality** of :
  - \* The **IF-OF** register
  - \* and the **Program Counter (PC)**
- \* To insert a **bubble**
  - \* Write a **bubble** (**nop** instruction) into the **OF-EX** register

# Data Path with Interlocks (Data-Lock)



# Ensuring the Branch-Lock Condition

- \* Option 1 :

- \* Use **delay slots** (**interlocks** not required)

- \* Option 2 :

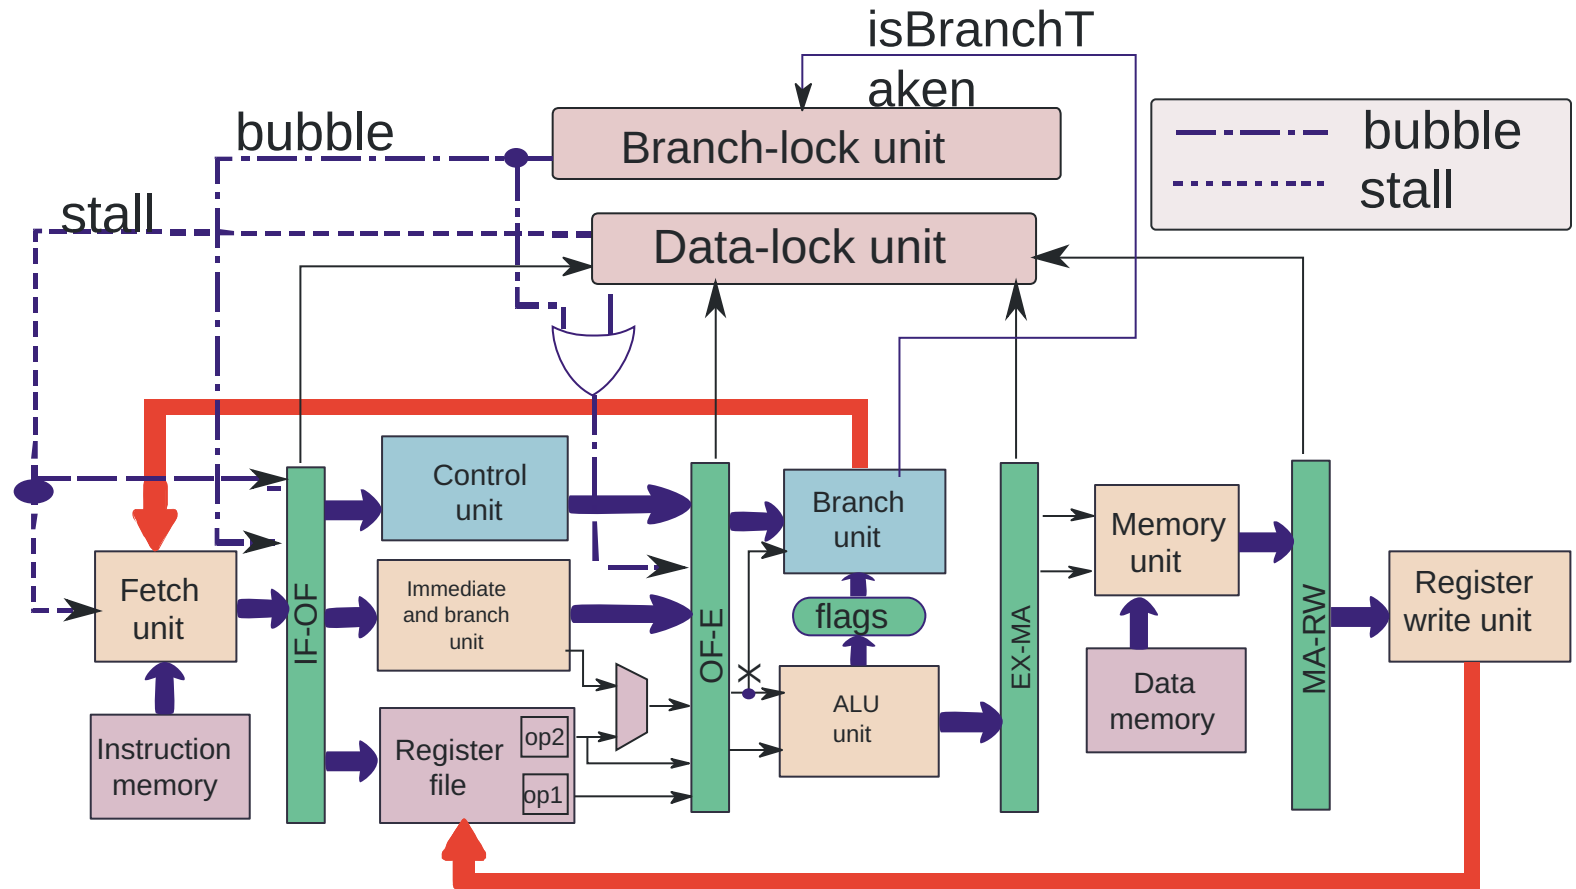
- \* **Convert** the instructions in the IF, and OF stages, to **bubbles** once a **branch instruction** reaches the **EX** stage.
  - \* Start fetching from the **next PC (not taken)** or the **branch target (taken)**

# Ensuring the Branch-Lock Condition - II

## \* Option 3


- \* If the **branch instruction** in the **EX** stage is **taken**, then **invalidate** the instructions in the IF and OF stages. Start **fetching** from the **branch target**.
- \* Otherwise, do not take **any special action**
- \* **This method is also called predict not-taken (we shall use this method because it is more efficient than option 2)**

# Data Path with Interlocks




# Outline

- \* Overview of Pipelining
- \* A Pipelined Data Path
- \* Pipeline Hazards
- \* Pipeline with Interlocks
- \* Forwarding
- \* Performance Metrics
- \* Interrupts/ Exceptions





# Relook at the Pipeline Diagram

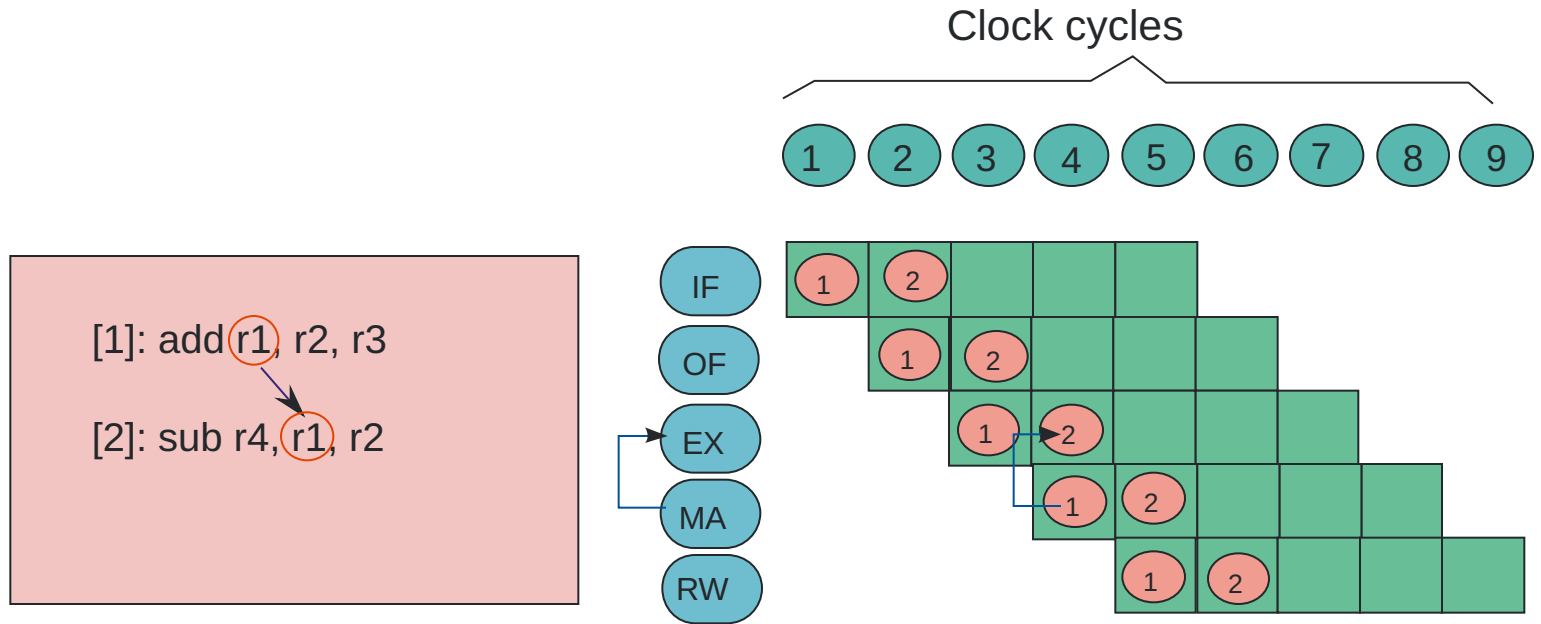


- \* (a) → with bubbles
- \* (b) → **no** bubbles (may lead to wrong results)

# Crucial Insight (Figure (b))

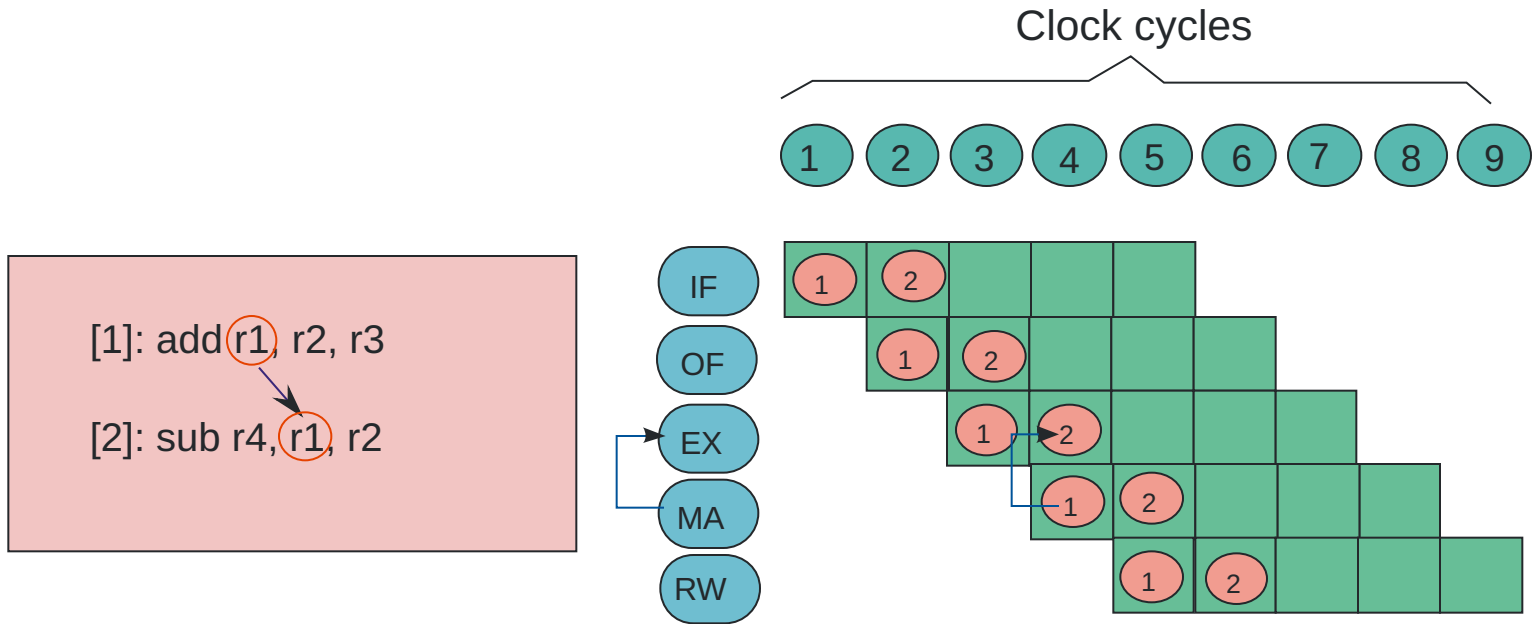
- \* When does instruction 2 need the value of r1 ?
  - \* **ANSWER** : Cycle 3, OF Stage (**wrong!!!**)
  - \* **CORRECT ANSWER** : Cycle 4, EX Stage
- \* When does instruction 1 produce the value of r1 ?
  - \* **ANSWER** : Cycle 5, RW Stage (**wrong!!!**)
  - \* **CORRECT ANSWER** : End of Cycle 3, EX Stage

# Forwarding



- \* If the **correct value** is already there in another stage, we can **forward** it.

# Forwarding from MA to EX



- \* Forwarding in cycle 4 from instruction [1] to [2]

# Different Forwarding Paths

- \* We need to add a multitude of forwarding paths
- \* Rules for creating forwarding paths
  - \* Add a path from a later stage to an earlier stage
  - \* Try to add a forwarding path as late as possible. For, example, we avoid the  $EX \rightarrow OF$  forwarding path, since we have the  $MA \rightarrow EX$  forwarding path
  - \* The IF stage is not a part of any forwarding path.

# Forwarding Path

- \* 3 Stage Paths

- \*  $RW \rightarrow OF$

- \* 2 Stage Paths

- \*  $RW \rightarrow EX$

- \*  $MA \rightarrow OF$  (**X Not Required**)

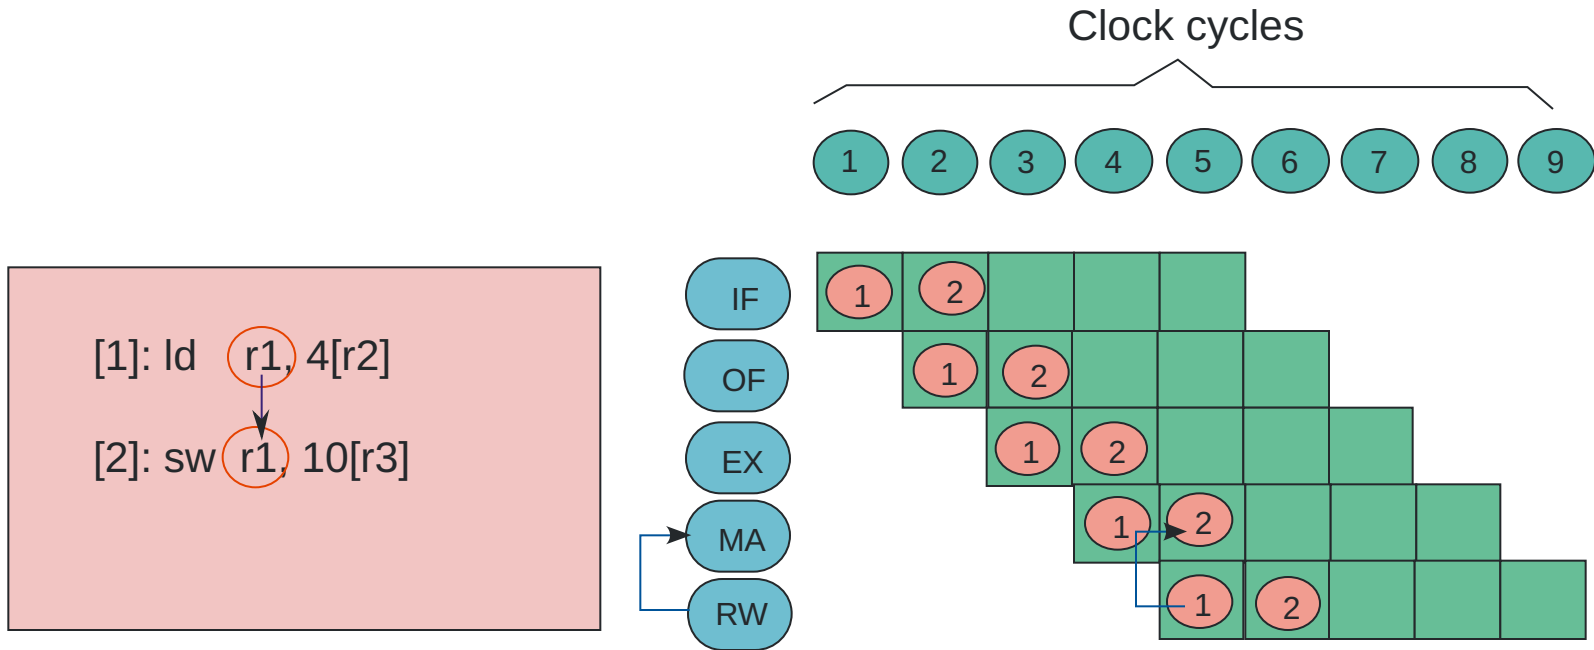
- \* 1 Stage Paths

- \*  $RW \rightarrow MA$  (load to store)

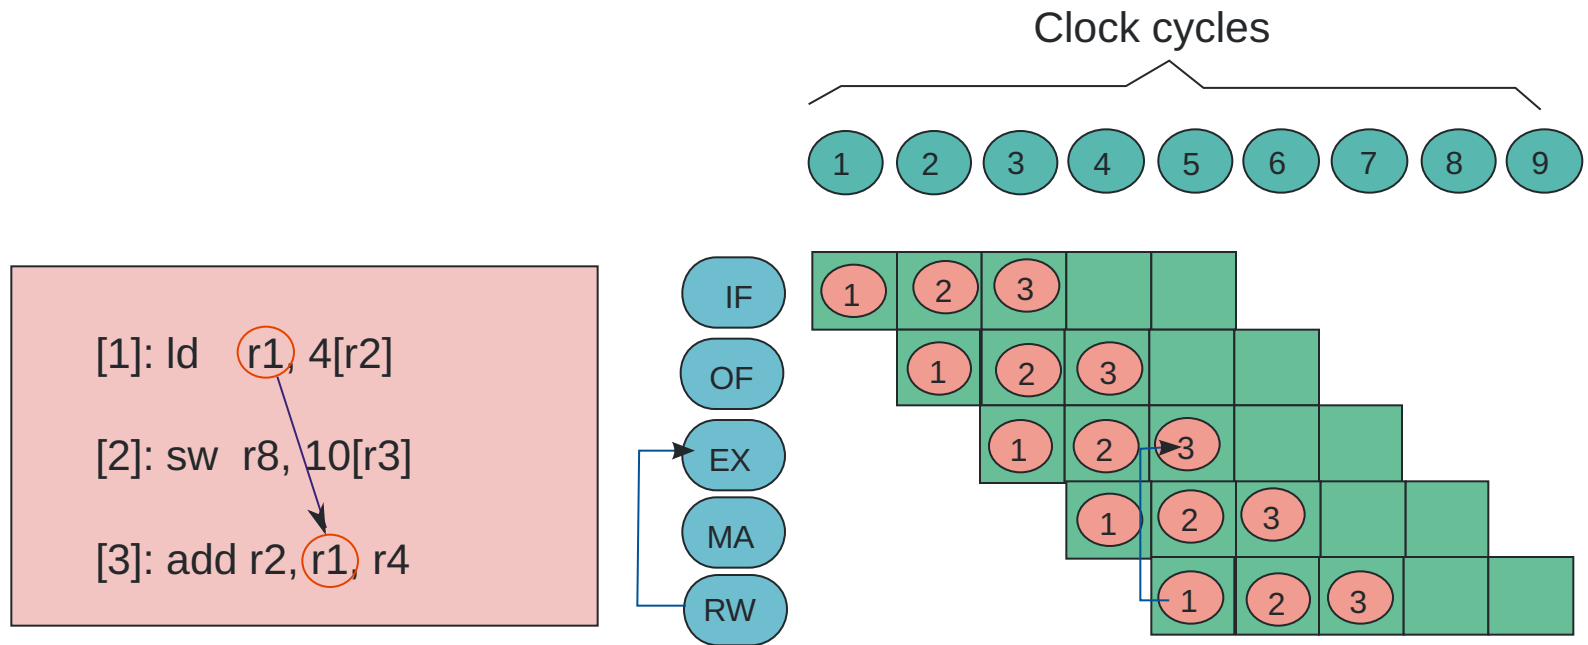
- \*  $MA \rightarrow EX$  (ALU Instructions, load, store)

- \*  $EX \rightarrow OF$  (**X Not Required**)

# Forwarding Paths : RW → MA

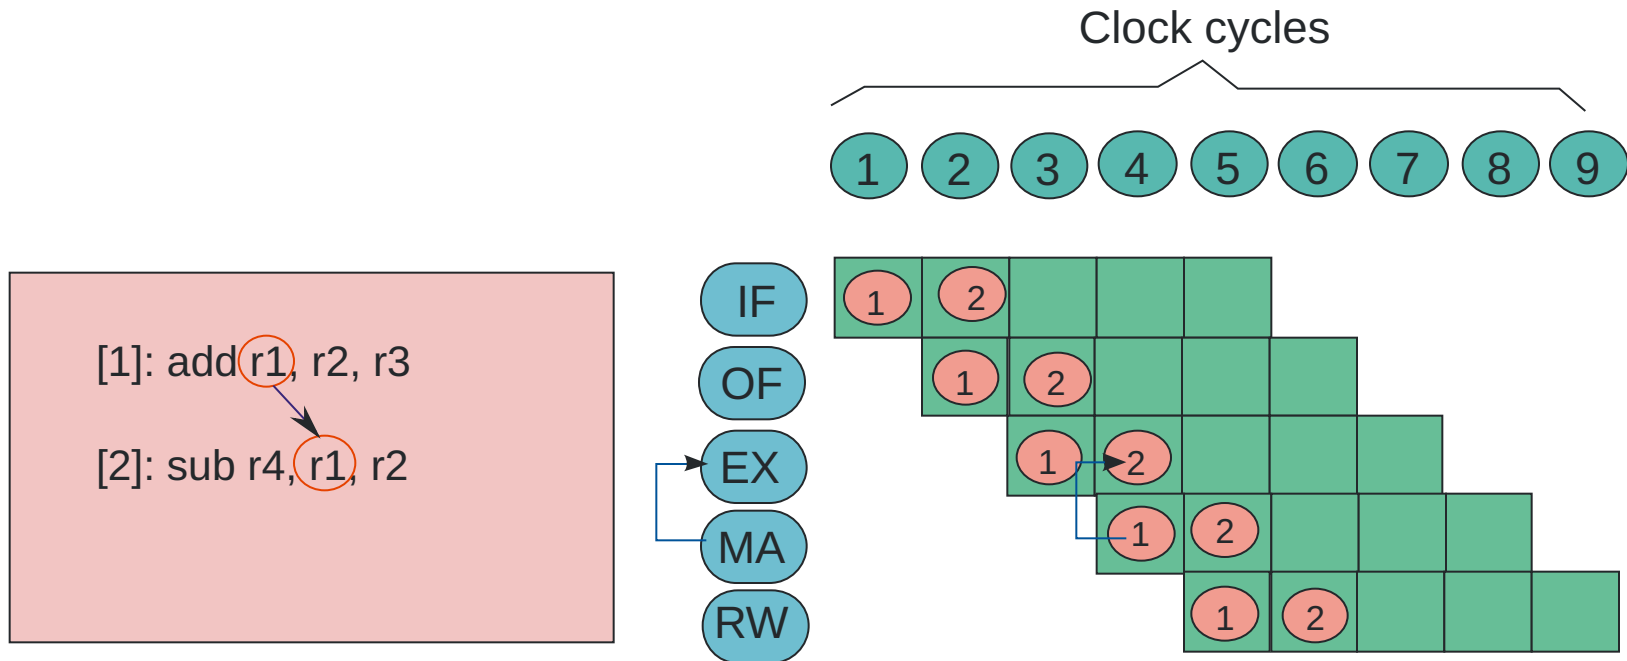


# Forwarding Paths : RW $\rightarrow$ EX

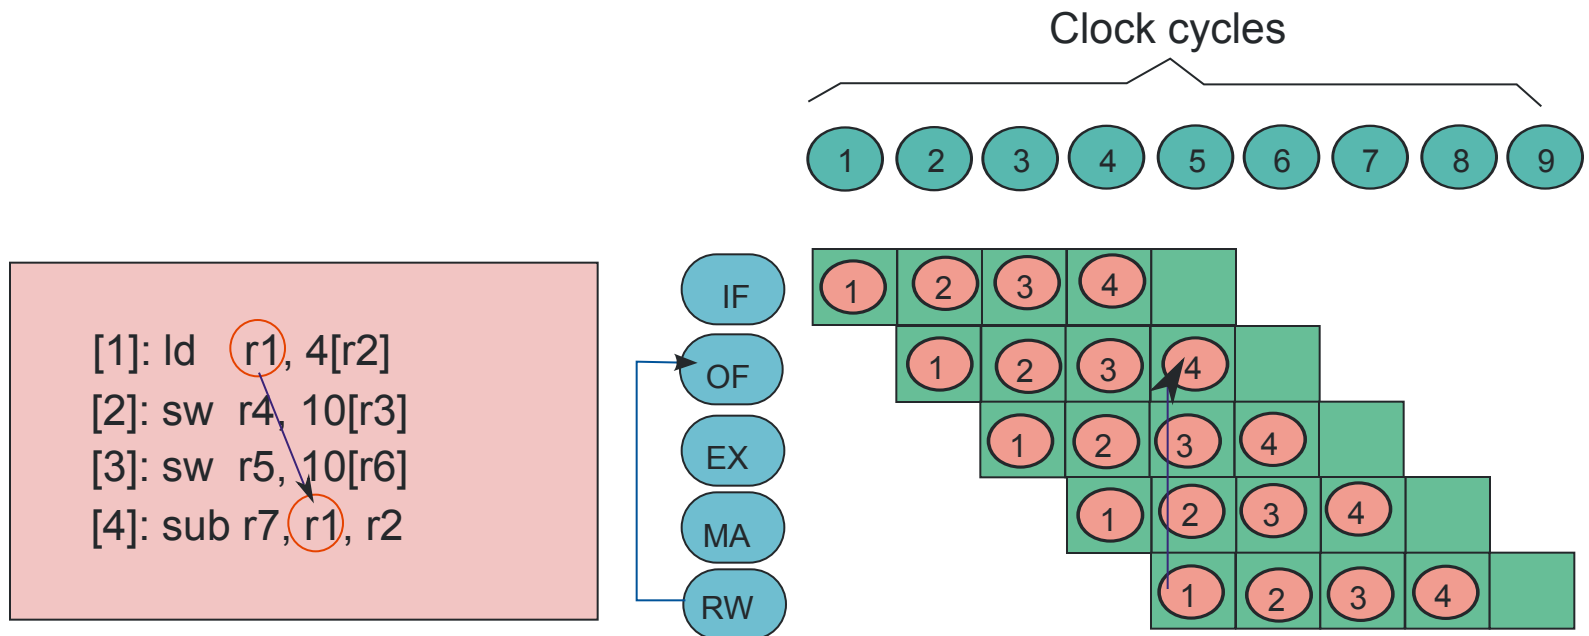




# Forwarding Path : MA $\rightarrow$ EX



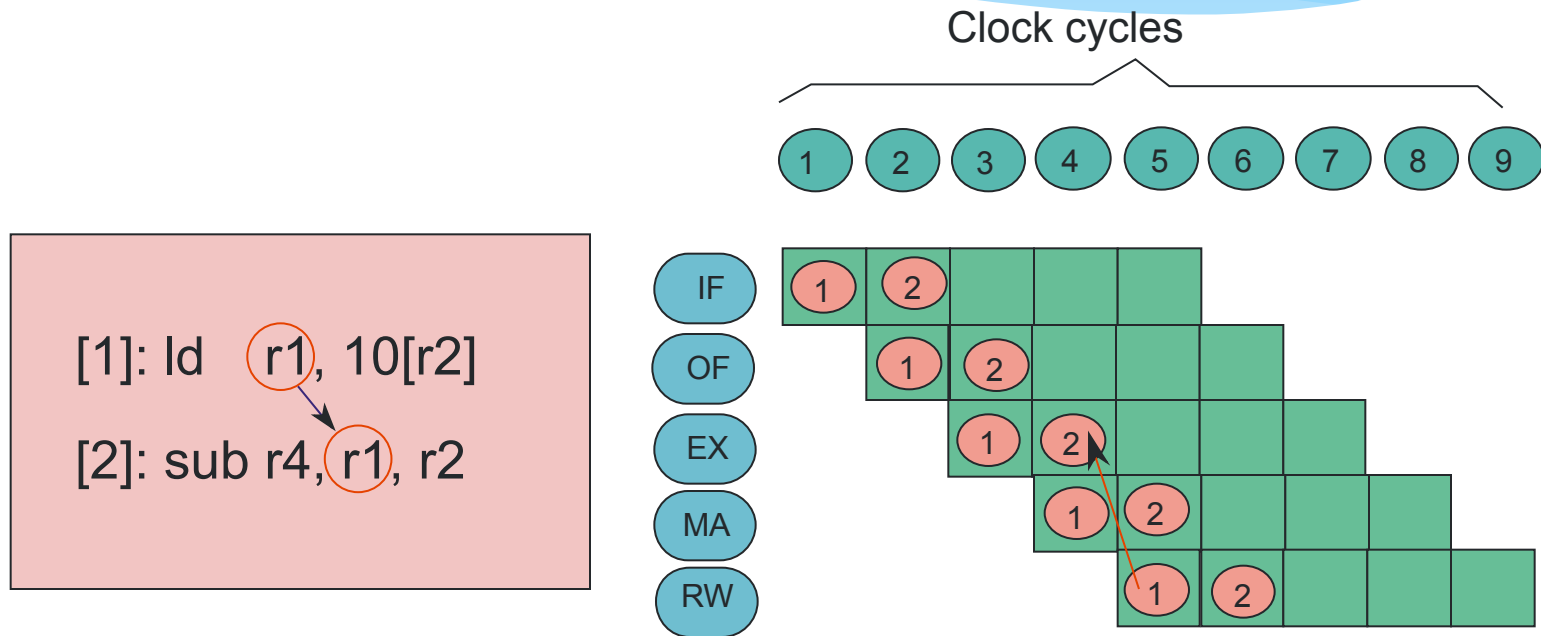
# Forwarding Path : RW → OF



# Data Hazards with Forwarding

- \* Forwarding has unfortunately not eliminated all data hazards
- \* We are left with one special case.
- \* Load-use hazard
  - \* The instruction immediately after a load instruction has a RAW dependence with it.

# Load-Use Hazard

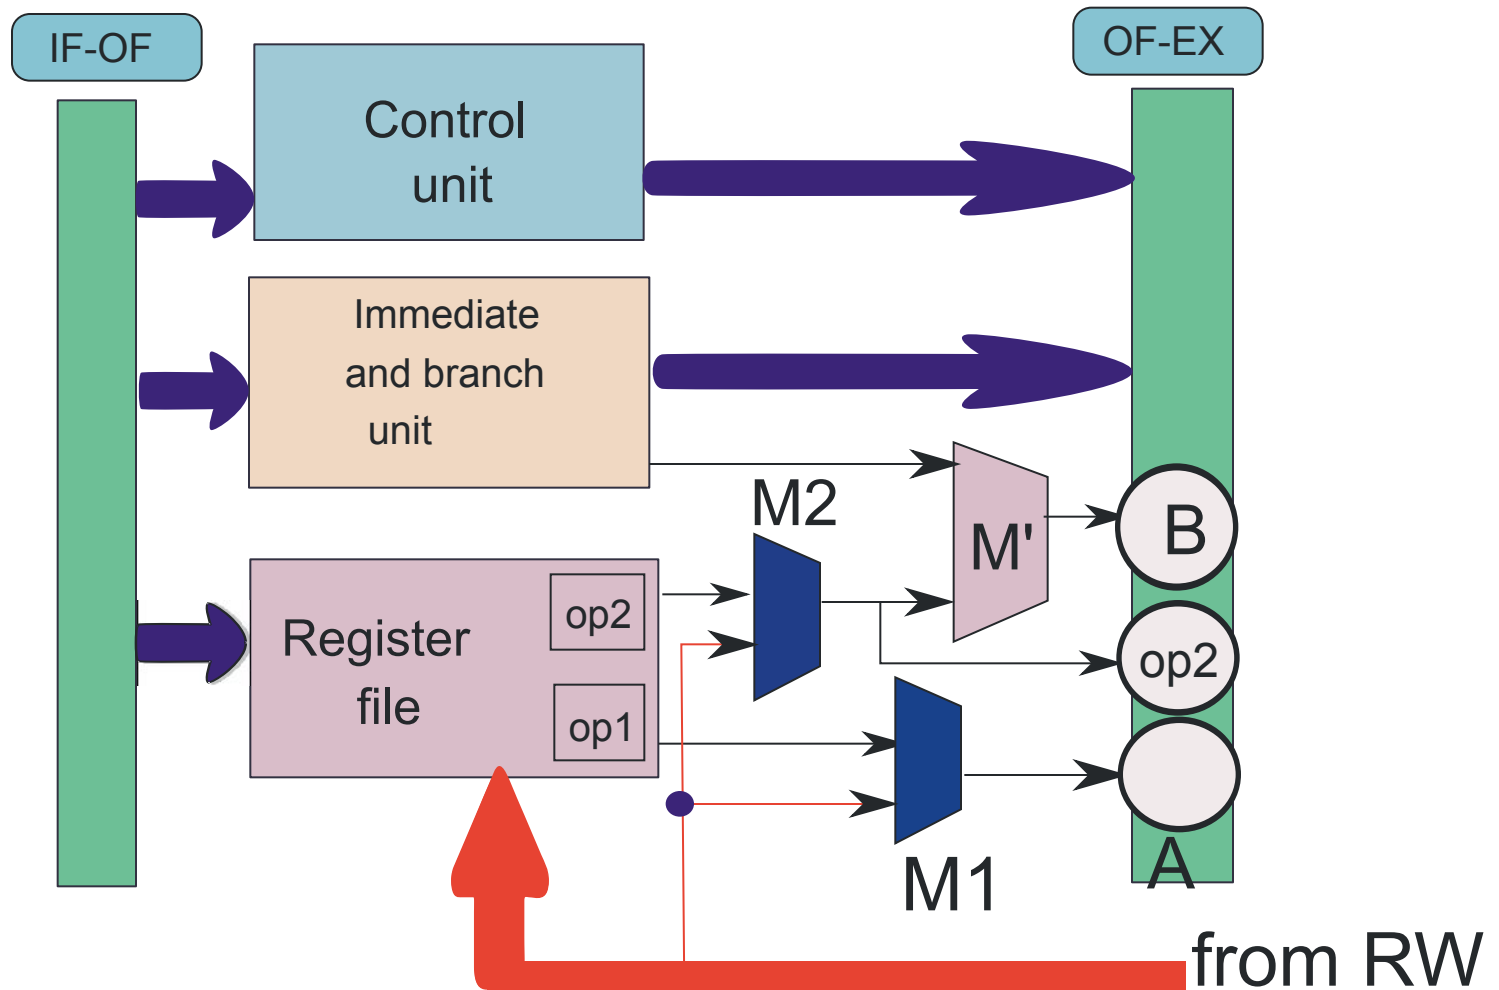


- \* **Cannot forward** (arrow goes backwards in time)
- \* Need to add a bubble (then use RW → EX forwarding)

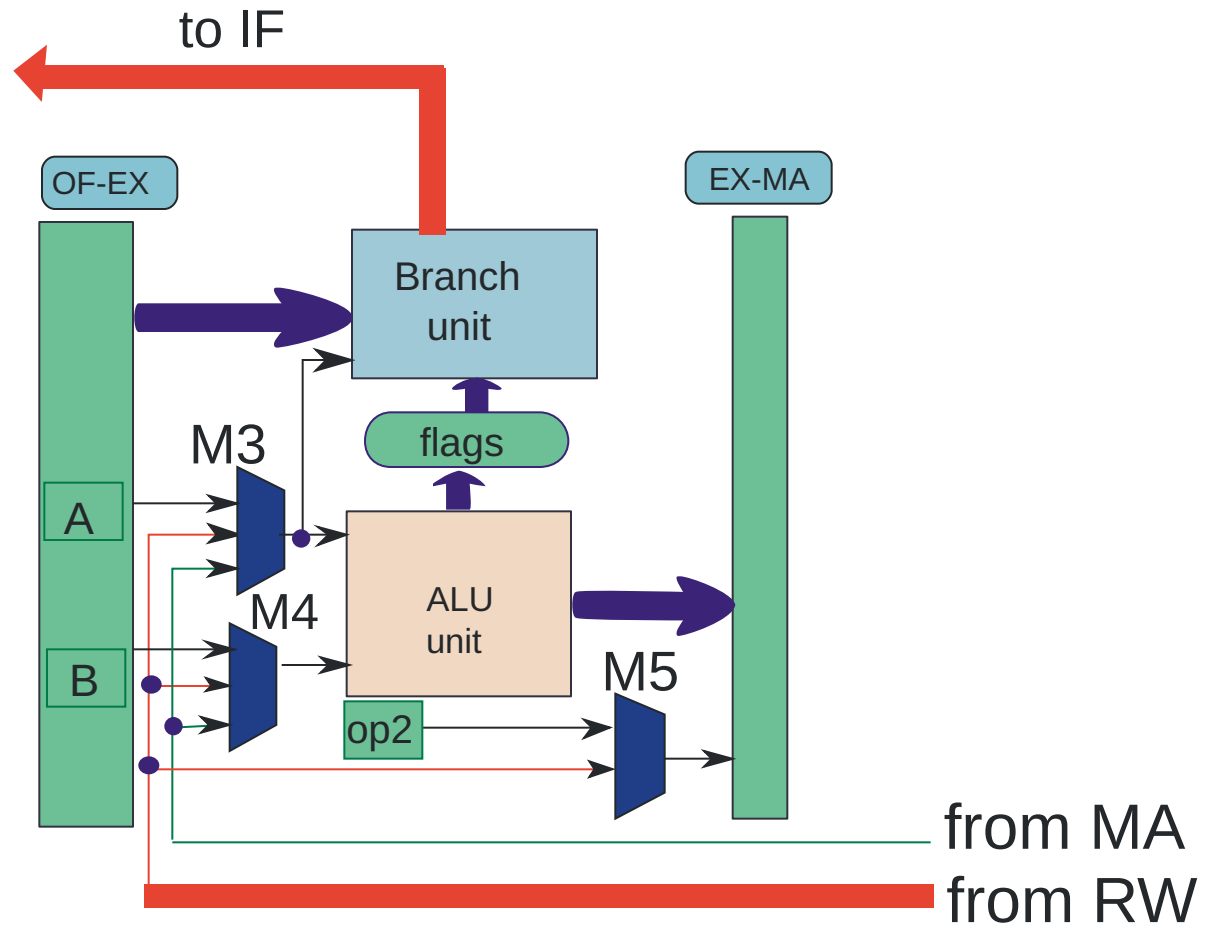
# Implementation of Forwarding

- \* At every stage there is a **choice** of inputs for each **functional unit**
  - \* Use the **default** inputs from the previous stage
  - \* **OR**, use one of the **forwarded** inputs
  - \* Use a **multiplexer** to choose between the inputs
  - \* A dedicated **forwarding unit** generates the signals for the **forwarding multiplexers**


# OF Stage with Forwarding



# EX Stage with Forwarding

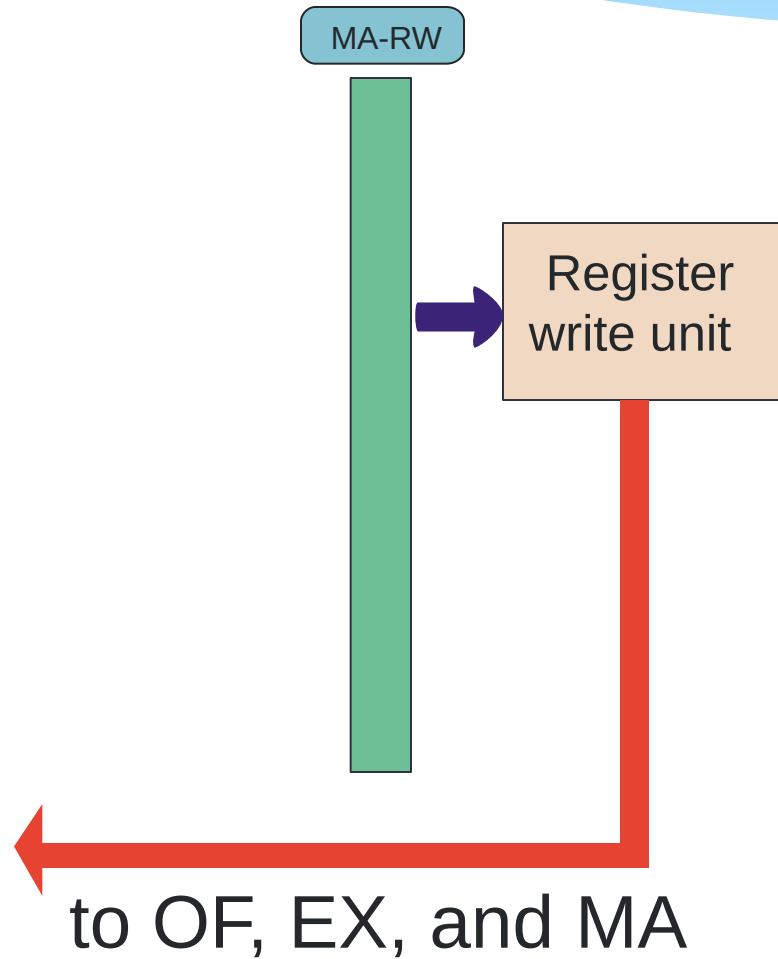


# MA Stage with Forwarding

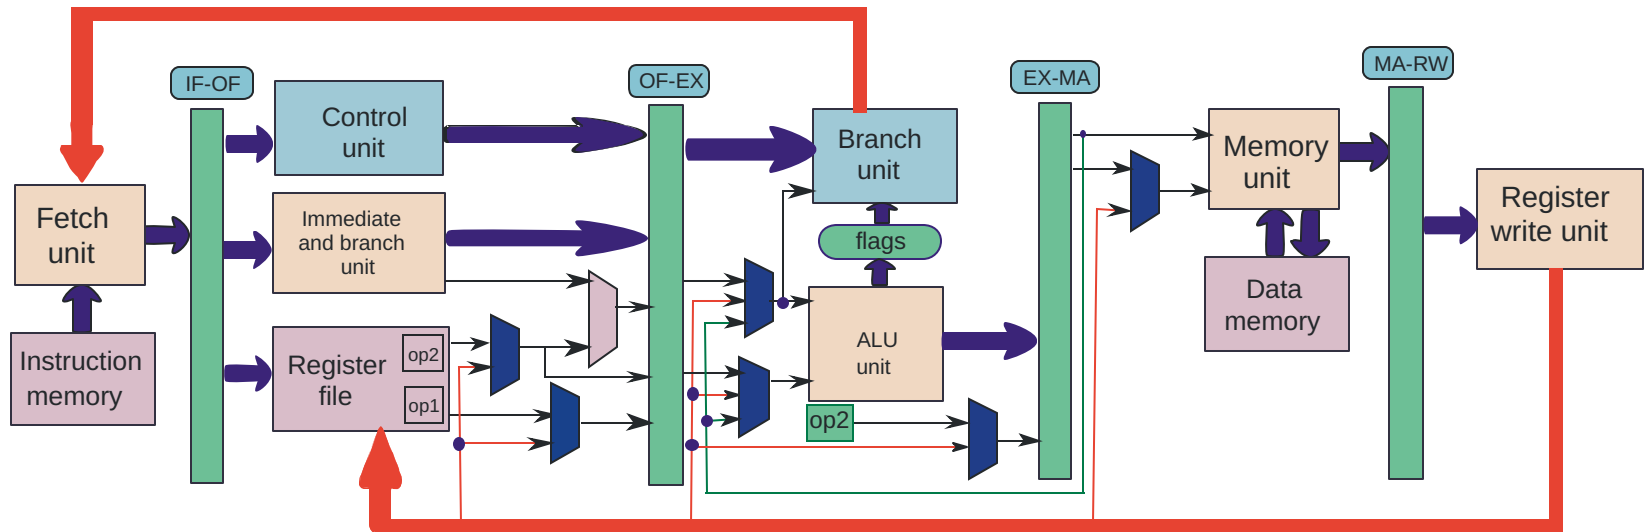




# RW Stage with Forwarding



# Data Path with Forwarding



# Forwarding Conditions

- \* Determine if there is a **conflict** between two **instructions** in different **stages**
  - \* Find if there is a **conflict** for the first operand (rs1/ ra)
  - \* Find if there is a **conflict** for the second operand (rs2/rd)
- \* Always **forward** from the **latest** instruction (that is earlier than the current instruction)

**Algorithm 6:** Conflict on the first operand (rs1/ra)

**Data:** instructions, [A], and [B] (possible forwarding: [B] → [A])

**Result:** conflict exists on rs1/ra (**true**), no conflict (**false**)

**if** [A].opcode  $\in$  (nop,b,beq,bgt,call,not,mov) **then**

    /\* Does not read from any register \*/

**return false**

**end**

**if** [B].opcode  $\in$  (nop, cmp, st, b, beq, bgt, ret) **then**

    /\* Does not write to any register \*/

**return false**

**end**

/\* Set the sources \*/

src1  $\leftarrow$  [A].rs1

**if** [A].opcode = ret **then**

    src1  $\leftarrow$  ra

**end**

/\* Set the destination \*/

dest  $\leftarrow$  [B].rd

**if** [B].opcode = call **then**

    dest  $\leftarrow$  ra

**end**

/\* Detect conflicts \*/

**if** src1 = dest **then**

**return true**

**end**

**return false**

**Algorithm 7:** Conflict on the second operand (rs2/rd)

**Data:** instructions, [A], and [B] (possible forwarding:  $[B] \rightarrow [A]$ )

**Result:** conflict exists on second operand (rs2/rd) (**true**), no conflict (**false**)

```
if [A].opcode  $\in$  (nop, b, beq, bgt, call) then
    /* Does not read from any register */
    return false
```

```
end
```

```
if [B].opcode  $\in$  (nop, cmp, st, b, beq, bgt, ret) then
    /* Does not write to any register */
    return false
```

```
end
```

```
/* Check the second operand to see if it is a register */
```

```
if [A].opcode  $\neq$  (st) then
    if [A].I = I then
        return false
    end
```


```
end
```

```
/* Set the sources */
```

```
src2  $\leftarrow$  [A].rs2
```

```
if [A].opcode = st then
    src2  $\leftarrow$  [A].rd
```

```
end
```



```
/* Set the destination */  
dest  $\leftarrow$  [B].rd  
if [B].opcode = call then  
    dest  $\leftarrow$  ra  
end  
/* Detect conflicts */  
if src2 = dest then  
    return true  
end  
return false
```

# Interlocks with Forwarding


## \* Data-Lock

- \* We need to only **check** for the **load-use hazard**
- \* If the **instruction** in the EX stage is a **load**, and the **instruction** in the OF stage uses its **loaded value**, then **stall** for 1 cycle

## \* Branch-Lock

- \* Remains the **same** as before.

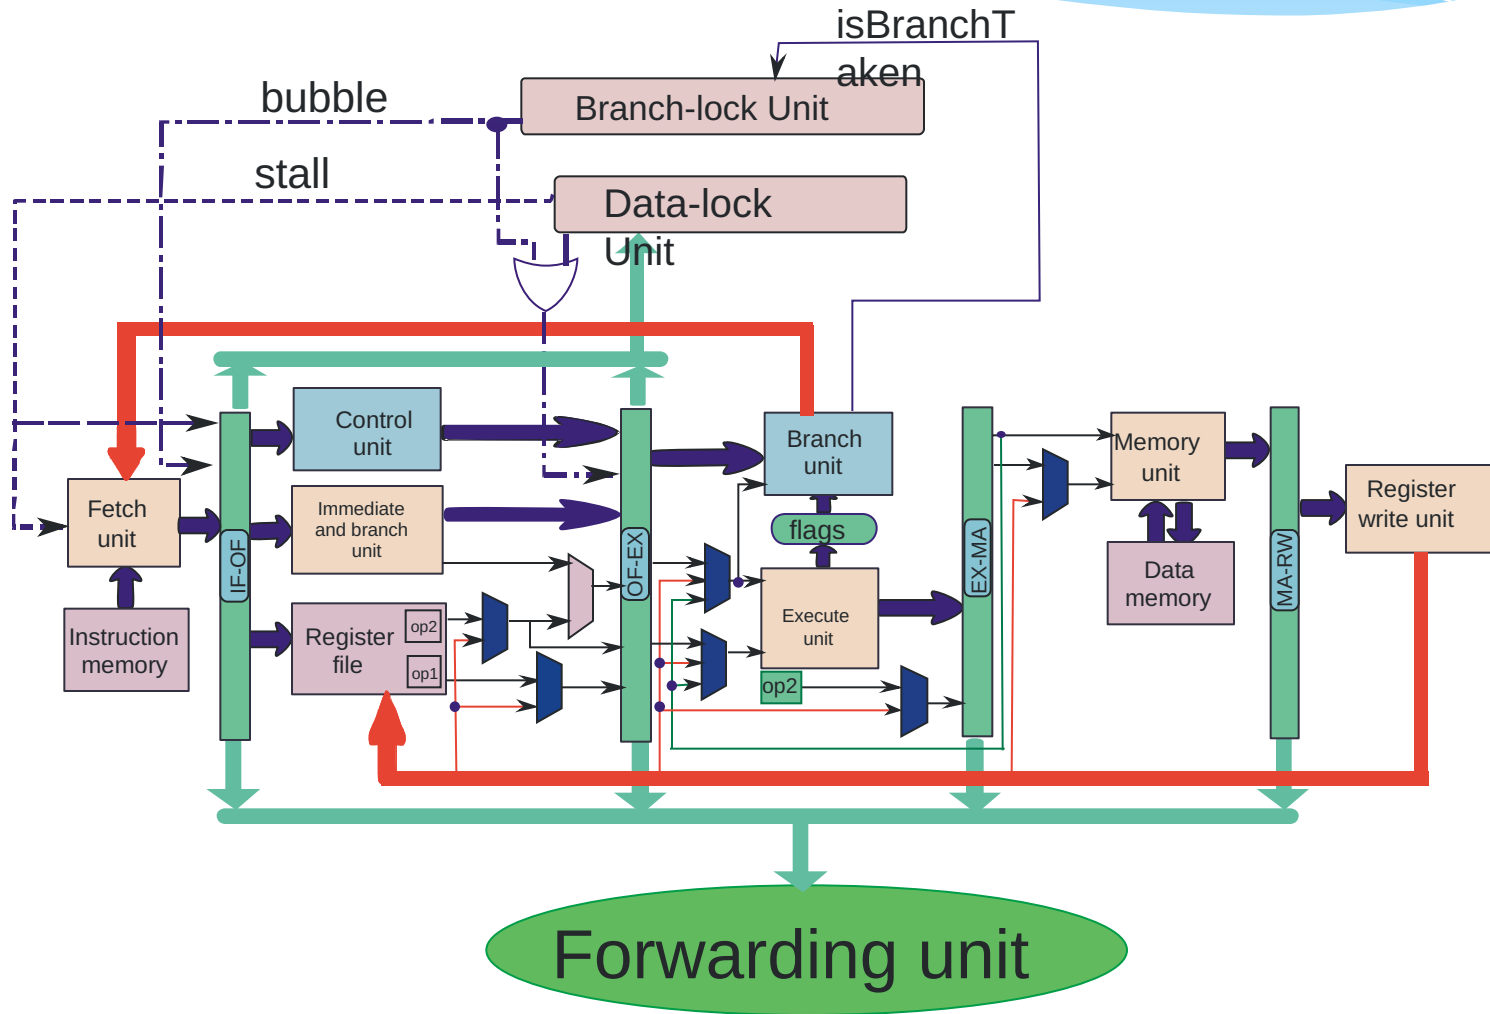
# The Curious Case of the *call* instruction



- \* The **call** instruction generates the **value** of **ra** in the RW stage
- \* Any instruction that uses the value of ra (such as ret) still works **correctly** !!!
- \* Prove it ....




# Complete Data Path



# Outline

- \* Overview of Pipelining
- \* A Pipelined Data Path
- \* Pipeline Hazards
- \* Pipeline with Interlocks
- \* Forwarding
- \* Performance Metrics
- \* Interrupts/ Exceptions



# Measuring Performance

- \* What do we mean by the **performance** of a processor ?
  - \* **ANSWER** : Almost nothing
- \* What should we ask instead ?
  - \* What is the **performance** with respect to a given program or a set of programs ?
  - \* **Performance** is **inversely proportional** to the time it takes to **execute** a **program**

# Computing the Time a Program Takes

$$\begin{aligned}\tau &= \#seconds \\ &= \frac{\#seconds}{\#cycles} * \frac{\#cycles}{\#instructions} * (\#instructions) \\ &= \underbrace{\frac{\#seconds}{\#cycles}}_{1/f} + \underbrace{\frac{\#cycles}{\#instructions}}_{CPI} * (\#instructions) \\ &= \frac{CPI * \#insts}{f}\end{aligned}$$

- \* **CPI** → **Cycles** per instruction
- \* **f** → **frequency** (cycles per second)

# The Performance Equation

$$P \propto \frac{IPC * f}{\#insts}$$

- \* **IPC** → 1/CPI (Instructions per Cycle)
- \* What are the **units** of performance ?
  - \* **ANSWER** : arbitrary

# Number of Instructions (#insts)

**Static Instruction:** The **binary** or executable of a **program**, contains a list of *static instructions*.

**Dynamic Instruction:** A **dynamic instruction** is a running instance of a static instruction, which is created by the **processor** when an instruction enters the pipeline.

- \* Note that these are **dynamic** instructions
  - \* **NOT** static instructions
- \* A smart compiler can reduce the number of executed instructions

# Number of Instructions(#insts) – 2

- \* Dead code removal

- \* Often **programmers** write **code** that does not determine the final output
- \* This code is **redundant**
- \* It can be identified and removed by the **compiler**

- \* Function inlining

- \* Very small **functions** have a lot of **overhead** → call, ret instructions, register spilling, and restoring
- \* Paste the code of the **callee** in the code of the **caller** (known as **inlining**)

# Computing the CPI

- CPI for a single cycle processor = 1
- CPI for an ideal pipeline (no hazards)
  - Assume we have  $n$  instructions, and  $k$  stages



# Computing the CPI

- \* CPI for a single cycle processor = 1
- \* CPI for an ideal pipeline(no hazards)
  - \* Assume we have  $n$  instructions, and  $k$  stages
  - \* The first instruction enters the pipeline in cycle 1
  - \* It leaves the pipeline in cycle  $k$
  - \* The rest of the  $(n-1)$  instructions leave in the next  $(n-1)$  consecutive cycles

$$CPI = \frac{n + k - 1}{n}$$

# Computing the Maximum Frequency

- \* Let the **maximum** amount of time that it takes to execute any **instruction** be :
  - \*  $t_{\max}$  (also known as **algorithmic work**)
- \* **Minimum** clock cycle time of a single cycle pipeline  $\rightarrow t_{\max}$
- \* In the case of a pipeline, let us assume that all the pipeline **stages** are **balanced**
- \* Time per stage  $\rightarrow t_{\max} / k$

# Maximum Frequency - II

- \* Let the **latch delay** be  $l$
- \* We thus have :

$$t_{stage} = \frac{t_{max}}{k} + l$$

$$\frac{1}{f} = \frac{t_{max}}{k} + l$$

The minimum cycle time ( $1/f$ ) is equal to  $t_{stage}$ . Let us thus, assume that our cycle time is as low as possible.

# Performance of an Ideal Pipeline

- \* Let us assume that the number of instructions are a **constant**

$$\begin{aligned} P &= \frac{f}{CPI} \\ &= \frac{\frac{t_{max}}{k} + l}{\frac{n + k - 1}{n}} \\ &= \frac{\left(\frac{t_{max}}{k} + l\right) * (n + k - 1)}{n} \\ &= \frac{(n - 1)t_{max}}{k} + (t_{max} + ln - l) + lk \end{aligned}$$

# Optimal Number of Pipeline Stages

$$\begin{aligned} \frac{\partial \left( \frac{(n-1)t_{max}}{k} + (t_{max} + l - l) + lk \right)}{\partial k} &= 0 \\ \Rightarrow -\frac{(n-1)t_{max}}{k^2} + l &= 0 \\ \Rightarrow k &= \sqrt{\frac{(n-1)t_{max}}{l}} \end{aligned}$$

- \*  $k$  is inversely proportional to  $\sqrt{l}$
- \*  $k$  is proportional to  $\sqrt{t_{max}}$

# Implications

- \* As we **increase** the **latch delay**, we should have **less** pipeline stages
  - \* We need to **minimise** the time wasted in accessing latches
- \* As we increase the **amount** of **algorithmic work**, we require more pipeline stages for ideal **performance**
  - \* More pipeline stages help **distribute** the work better, and increase the **overlap** across instructions

# A Non-Ideal Pipeline

- \* Our ideal CPI ( $CPI_{ideal} = 1$ ) is 1
- \* However, in reality, we have stalls

$$CPI = CPI_{ideal} + stall\_rate * stall\_penalty$$

- \* Let us assume that the **stall rate** is a function of the **program**, and its nature of **dependences**

# Non-Ideal Pipeline - II

- \* Let us assume that the **stall penalty** is **proportional** to the number of **pipeline stages**
- \* Both these assumptions are strictly **not correct**. They are being used to make a coarse grained **mathematical model**.
- \*  $CPI = (n+k-1)/n + rck$ 
  - \*  $r \rightarrow$  **stall rate**,  $c \rightarrow$  constant of **proportionality**



# Mathematical Model

$$\begin{aligned}
 P &= \frac{f}{CPI} \\
 &= \frac{\frac{1}{\frac{t_{max}}{k} + l}}{\frac{n+k-1}{n} + rck} \\
 &= \frac{n}{\frac{(n-1)t_{max}}{k} + (rcnt_{max} + t_{max} + ln - l) + lk(1 + rcn)}
 \end{aligned}$$

# Mathematical Model - II

$$\frac{\partial \left( \frac{(n-1)t_{max}}{k} + (rcnt_{max} + t_{max} + \ln - l) + lk(1 + rcn) \right)}{\partial k} = 0$$

$$\Rightarrow -\frac{(n-1)t_{max}}{k^2} + l(1 + rcn) = 0$$

$$\Rightarrow k = \sqrt{\frac{(n-1)t_{max}}{l(1+rcn)}} \approx \sqrt{\frac{t_{max}}{lrc}} \quad (as \ n \rightarrow \infty)$$

# Implications

- \* For programs with a lot of **dependences** (high value of **r**) → Use **less pipeline** stages
- \* For a pipeline with **forwarding** → **c** is **smaller** (than a **pipeline** that just has interlocks)
  - \* It requires a larger number of **pipeline stages** for optimal performance

# Implications

- \* The optimal number of pipeline stages is directly proportional to  $\sqrt{(t_{\max} / I)}$ 
  - \* This ratio is not significantly changing across technologies.
  - \* This explains why the number of pipeline stages has remained more or less constant for the last 5-10 years

# Example

**Example** Consider two programs that have the following characteristics.

| Program 1             |          | Program 2             |          |
|-----------------------|----------|-----------------------|----------|
| Instruction Type      | Fraction | Instruction Type      | Fraction |
| loads                 | 0.4      | loads                 | 0.3      |
| Branches              | 0.2      | Branches              | 0.1      |
| ratio(taken branches) | 0.5      | ratio(taken branches) | 0.4      |

*The ideal CPI is 1 for both the programs. Let 50% of the load instructions suffer from a load use hazard. Assume that the frequency of  $P_1$  is 1, and the frequency of  $P_2$  is 1.5. Here, the units of the frequency are not relevant. Compare the performance of  $P_1$  and  $P_2$ .*

# Example

**Answer:**

$$\begin{aligned} CPI_{new} = & CPI_{ideal} + 0.5 \times (\text{ratio}(\text{loads})) \times 1 \\ & + \text{ratio}(\text{branches}) \times \text{ratio}(\text{taken branches}) \times 2 \end{aligned} \quad (9.13)$$


*We thus have:*

$$CPI_{P_1} = 1 + 0.5 \times 0.4 + 0.2 \times 0.5 \times 2 = 1 + 0.2 + 0.2 = 1.4$$

$$CPI_{P_2} = 1 + 0.5 \times 0.3 + 0.1 \times 0.4 \times 2 = 1 + 0.15 + 0.08 = 1.23$$


*The performance of  $P_1$  can be expressed as  $f/CPI = 1 / 1.4 = 0.71$  (arbitrary units). Similarly, the performance of  $P_2$  is equal to  $f/CPI = 1.5/1.23 = 1.22$  (arbitrary units). Hence,  $P_2$  is faster than  $P_1$ . We shall often use the term, arbitrary units, a.u., when the choice of units is irrelevant.*

# Performance, Architecture, Compiler



# Outline

- \* Overview of Pipelining
- \* A Pipelined Data Path
- \* Pipeline Hazards
- \* Pipeline with Interlocks
- \* Forwarding
- \* Performance Metrics
- \* Interrupts/ Exceptions





# What happens when you press a key ?

- \* The **keyboard** logs the key press
  - \* **Converts** the key to **ASCII** or Unicode
  - \* **Sends** the code to the **processor**
  - \* The **processor** thus receives an **interrupt**
  - \* It **suspends** the current **program**
  - \* Jumps to the **interrupt handler**.
  - \* The **interrupt handler** draws the shape associated with the key
  - \* The **processor** returns to execute the original **program**

# Exceptions

- \* **Exceptions** are generated when
  - \* A **program** accesses an illegal address
  - \* We try to **divide** 5/0
  - \* We issue an **invalid instruction**
  - \* ...
- \* **Exception** are treated the same way as **interrupts**
  - \* Jump to the **exception** handler
  - \* Come back and start executing **programs**

# Precise Exceptions

## \* Informal definition

- \* We need to **return** to the original **program** at exactly the same **point**, at which we had left it
- \* The execution of the **interrupt handler** should not disrupt the execution of the original **program** in any way. The outcome of the original program should be independent of the **interrupt** (unless the program caused an **exception**).

# Precise Exceptions - II

- \* **Formal Definition**

- \* Let us number the **dynamic instructions** in a program :  
 $I_1 \dots I_n$
- \* Let us assume that an instruction **completes** after it either updates memory, writes to registers, or reaches the MA stage (cmp, b, beq, bgt, ret)
- \* Let the last **program instruction** that **completes** before the **first** instruction in the **interrupt handler completes**, be  $I_k$
- \* Let all the **program** instructions that **complete** before the first instruction in the interrupt handler **completes**, be **C**

# Precise Exceptions - III

$$I_j \in C \Leftrightarrow (j \leq k)$$

- \* We need to ensure this **condition**.
- \* Also no **instruction** of the form  $I_{k'}$  ( $k' > k$ ) should **complete** before all the instructions in the **interrupt handler complete**
- \* After returning, we can seamlessly **execute**  $I_k$  (same instruction) or  $I_{k+1}$  (next instruction)

# Marking Instructions

- \* When an **interrupt** arrives, let us **mark** the instruction in the **MA** stage
- \* Otherwise, if there is an exception
  - \* We mark the instruction, as soon as it encounters a fault/exception
- \* Once, an instruction is **marked**, we have two kinds of **instructions** in the pipeline
  - \* Instructions **before/after** the **marked** instruction

# Implementing Precise Exceptions

- \* Wait till a **marked instruction** reaches the end of the **pipeline**
  - \* Convert all the **instructions** after the marked instruction to **bubbles**
- \* Ensures that both the **conditions** of a **precise exception** are met.
- \* Once the **marked** instruction reaches the end of the **pipeline**, the exception unit loads the **pc** of the **interrupt handler**

# Saving/Restoring Program State

- \* Program State

- \* PC
- \* Registers
- \* Flags
- \* Memory

- \* **Memory** → Assume that there is no **overlap** of memory regions, unless explicitly intended



# oldPC Register

- \* Let us add a **npc** field in the instruction packet
  - \* For taken branches it is equal to the branch target
  - \* For all other instructions it is equal to  $(pc + 4)$
- \* We populate the **npc** field in the EX stage
- \* Depending on the type of the exception, we might want to return to **pc** or **npc**
  - \* The **exception unit** sets the **oldPC** register to the right **return address**

# Spilling/ Restoring Registers

- \* In the case of **functions**, we stored registers on the **stack**
  - \* In this case, the **interrupt handler** has a separate stack.
  - \* We cannot **overwrite** the stack pointer (we will lose its previous value)
- \* **Solution :**
  - \* Use an additional register, **oldSP**, to save the stack pointer of the program.
  - \* Load the new stack pointer, and spill all the registers
  - \* Save oldPC

# The Strange Case of the Flags

## \* Naive solution :

- \* Do not allow any instruction after the marked instruction to update the flags register
- \* We detect an exception typically towards the middle or end of a cycle
- \* By that time, the instruction might have already updated the flags register (at least the master latch)

# Solution

- \* Add a **flags** field to the **instruction packet**
- \* Every **instruction** saves the updated value of the **flags register** to the **flags** field in its instruction packet.
- \* The exception unit saves the flags field of the marked instruction to the **oldFlags** register
- \* We store the **oldFlags** register to a location in the interrupt handler's **stack**

# Privileged Instructions

- \* We add the following new **registers** to the ISA
  - \* **special registers** → flags, oldFlags, oldPC, oldSP
  - \* The **flags** register is now accessible to the ISA
- \* Add the **z series** of privileged instructions
- \* **movz** instruction
  - \* Transfers values between **regular** registers and **special** registers

# Privileged Instructions - II

- \* **retz**

- \*  $pc \leftarrow \text{oldPC}$

- \* The **z series** instructions can only be used by the interrupt handler, and the operating system

- \* We use a **CPL** bit (current privilege level) bit for setting **permissions**

- \* User programs (CPL = 1)

- \* Interrupt handlers, kernel programs, (CPL = 0)

# Implementing *movz* and *retz*

- \* Define two new opcodes
  - \* These *opcodes* are usable only when  $CPL = 0$
  - \* Augment the OF and RW stages to use the *special registers*.
  - \* *movz*, and *retz* see a different view of registers.

| Register        | Encoding |
|-----------------|----------|
| <i>r0</i>       | 0000     |
| <i>oldPC</i>    | 0001     |
| <i>oldSP</i>    | 0010     |
| <i>flags</i>    | 0011     |
| <i>oldFlags</i> | 0100     |
| <i>sp</i>       | 1110     |

# Assembly Code for Spilling Registers

```
/* save the stack pointer */
movz oldSP, sp
mov sp, 0x FF FC

/* spill all the registers other
than sp*/
st r0, -4[sp]
st r1, -8[sp]
st r2, -12[sp]
st r3, -16[sp]
st r4, -20[sp]
st r5, -24[sp]
st r6, -28[sp]
st r7, -32[sp]
st r8, -36[sp]
st r9, -40[sp]
st r10, -44[sp]
st r11, -48[sp]
st r12, -52[sp]
```



# Spilling Registers - II

```
st r13, -56[sp]
st r15, -60[sp]

/* save the stack pointer */
movz r0, oldSP
st r0, -64[sp]

/* save the flags register */
movz r0, oldFlags
st r0, -68[sp]

/* save the oldPC */
movz r0, oldPC
st r0, -72[sp]

/* update the stack pointer */
sub sp, sp, 72
/* code of the interrupt handler */
....
....
....
```

# Restoring Registers

```
/* update the stack pointer */
add sp, sp, 72

/* restore the oldPC register */
ld r0, -72[sp]
movz oldPC, r0

/* restore the flags register */
ld r0, -68[sp]
movz flags, r0

/* restore all the registers other than sp*/
ld r0, -4[sp]
ld r1, -8[sp]
ld r2, -12[sp]
ld r3, -16[sp]
ld r4, -20[sp]
ld r5, -24[sp]
ld r6, -28[sp]
ld r7, -32[sp]
ld r8, -36[sp]
```

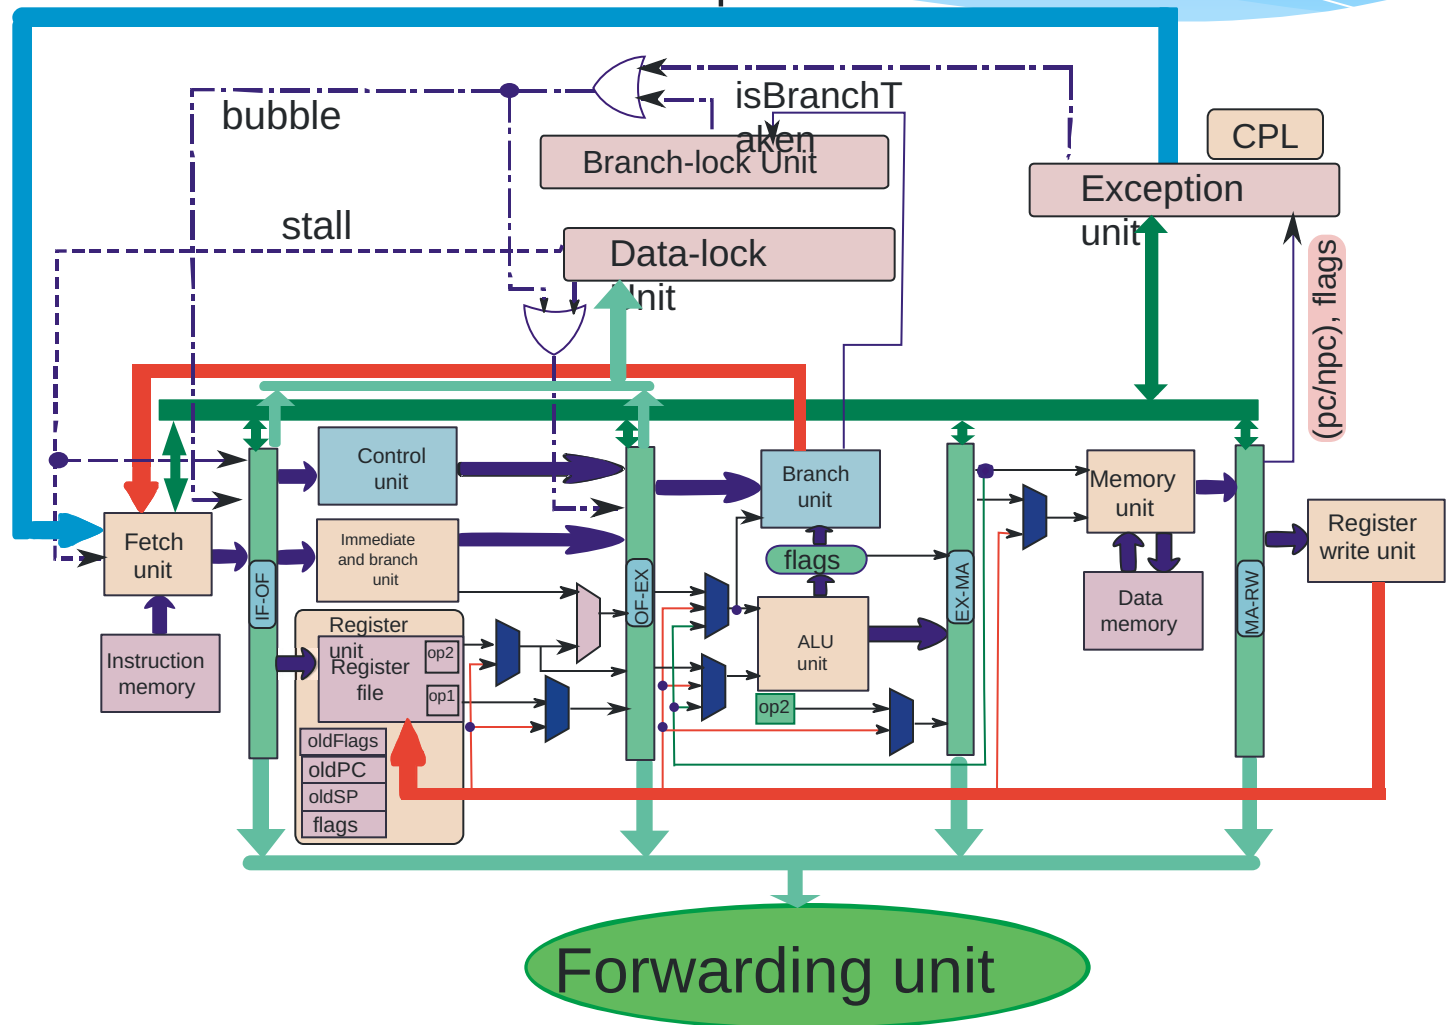
# Restoring Registers – II

```
ld r9, -40[sp]
ld r10, -44[sp]
ld r11, -48[sp]
ld r12, -52[sp]
ld r13, -56[sp]
ld r15, -60[sp]

/* restore the stack pointer */
ld sp, -64[sp]

/* return to the program */
retz
```

## PC of exception handler





# THE END